EGOISM

Jonathan Wolff

Ask a man why he uses exercise*; he will answer,* because he desires to keep his health. *If you then enquire,* why he desires health*, he will readily reply,* because sickness is painful. *If you push your enquiries farther, and desire a reason,* why he hates pain*, it is impossible he can ever give any. This is an ultimate end, and is never referred to any other object*.

—David Hume, An Enquiry Concerning the Principles of Morals

WHY BE MORAL?

One of the earliest systematic discussions of moral philosophy in the Western philosophical tradition occurs in the Ancient Greek philosopher Plato's great masterpiece, *The Republic* (c. 380 bce). In that text, Plato defends his own view of the nature of morality … However, in an argumentative strategy that we all can learn from, Plato is not content with brushing aside his opponent's view. Instead he sets out the opposing view in its strongest possible form. We can understand why he does this: If he can find the strongest version of his opponent's view, and still refute it, then his own view is firmly defended.

Plato therefore tells us a story about someone who is able to get away with acting immorally but with no adverse consequences. This is the mythical story of the "Ring of Gyges." According to the story, Gyges was a shepherd who found a magic ring that, if twisted around, made its wearer invisible. Suppose you had such a ring. What would you do? In Plato's *Republic*, a character named Glaucon represents the skeptical view in a dialogue with Socrates; he suggests the following:

> No man is so unyielding that he would remain obedient to justice and keep his hands off what does not belong to him if he could steal with impunity in the very midst of the public market itself. The same if he could enter into houses and lie with whom he chose, or if he could slay—or release from bondage—whom he would, behaving toward other men in these and all other things as if he were the equal of a god. The just man would act no differently from the unjust; both would pursue the same course.

One might argue that here is the great proof that no one is willingly just; men will be just only if constrained. This is because every man believes that justice is not really to his interest. If he has the power to do wrong, he will do wrong, for every man believes in his heart that injustice will profit him more than justice.

These are the settled convictions of all those who choose to adopt them. They hold that anyone who acquires extraordinary power and then refuses to do wrong and plunder others is truly to be pitied (and a great fool as well). Publicly, however, they praise the fool's example, convinced that they must deny what they really think so that they will

not encourage unjust acts against themselves. (1999, p. 56)

Who could resist using the Ring of Gyges? I think a lot of people would be tempted to use their new power to do such things as looking around the homes of the rich and famous; snooping and listening in on conversations, especially conversations about themselves; and doing other disreputable, slightly thrilling, things. But how far would you go? Steal? Kill? If not, why not? Is it because morality does have strong independent force after all? Or is it because we have been brought up not to harm others, and these rules are difficult to break whether or not we think they are ultimately justified? But perhaps, once you really were convinced that you could not be caught, you would kill if that were the only way of getting what you wanted. Or to put a rather different question: Would it be rational to stick with the moral rules if you knew that you could get away with breaking them?

Reflecting on this example suggests at least two different theories we need to consider. The first is what we could call a psychological claim, one about human behavior. This view suggests, as Glaucon presents it, that human beings cannot help pursuing their own self-interest. This theory has become known as **psychological egoism**; and if it is true, it would rule out the possibility of acting morally unless that behavior happened to coincide with self-interest. The second is a moral claim: Human beings have the right, perhaps even the duty, to pursue their own self-interest to the exclusion of the interests of others. For obvious reasons, this theory is often called **ethical egoism**. Although closely related, these theories are not the same. For example, some ethical egoists believe we are capable of acting altruistically, but we would be morally wrong to do so (strange though that may sound). Therefore, they deny psychological egoism. It is, though, possible to hold both views; and both have serious implications for morality and for moral philosophy. Let's turn first to psychological egoism.

PSYCHOLOGICAL EGOISM

Ultimately, in the psychological egoist's view, everything we do is aimed at making things better for ourselves. We are destined to seek pleasure, happiness, or even feelings of self-worth. Of course, we are often rather clever about it, making it look like we care about others or would sacrifice our own interests for the common good. But deep down, in this view it is always the same. All we do—all we can do—is follow our self-interest. But if that is true, then it seems that morality, properly speaking, is squeezed out from the start. For if you have no choice but to follow your self-interest, how can you also act morally? Surely morality, from time to time, requires genuine self-sacrifice.

This skeptical point about morality can be set out as consisting of an argument from two premises to a conclusion:

> Premise 1: Human beings are constituted by nature to pursue their own self-interest.

> Premise 2: Morality often requires self-denial (i.e., acting against your self-interest).
>
> Therefore
>
> Conclusion: Human beings are not capable of acting morally.

Those who are impressed by this argument sometimes draw the further conclusion that morality is somehow redundant or perhaps even deceptive—a myth we tell ourselves, perhaps to feel better about our selfishness.

Let's examine the logic of this simple argument more closely. The critic of morality holds premise 1, the psychological egoist claim that humans are constituted to follow our own self-interest, and premise 2—that morality often requires self-denial—and claims it follows that human beings cannot act morally. Is the critic right? One question is whether the conclusion really does follow from the premises. Is it a **logically valid** argument? If we look at it again, can we poke holes? Being maximally critical, is it possible to think of a situation in which the premises are true but the conclusion false? As we saw in Chapter 1, this is how philosophers normally attempt to determine whether an argument is logically valid: They test it to the point of attempted destruction. And in this case we can in fact question the argument. Look at the conclusion: Human beings are not capable of acting morally. Why is that? Because, according to the premises, we are constituted to follow our self-interest; and morality often requires self-denial. But does morality *always* require self-denial? Perhaps, but that's not what premise 2 says. It says that morality *often* requires self-denial. All that follows, then, is that *often* human beings are not capable of acting morally. And so we can see that the conclusion is ambiguous between

> Conclusion a (modest version): Human beings are *often not capable* of acting morally.
>
> Conclusion b (strong version): Human beings are *never capable* of acting morally.

The modest version of the conclusion does seem to follow logically from the premises, but the stronger version does not. That doesn't mean that it is false—it may be that human beings are never capable of acting morally. The point is that the argument we have been given so far does not establish such a claim. The argument allows that where morality is consistent with self-interest, we can in fact act morally. Suppose, for example, you have promised to look after your neighbor's dog for the day. So now you have a moral obligation to do so. But suppose looking after this particular dog is lots of fun and one of your favorite things to do. In this case morality and self-interest seem to coincide. Therefore, acting self-interestedly can be compatible with morality even if they can also clash. For example, you might also have promised to look after your neighbor's cat, which you view as a nuisance.

We can, though, push the discussion to a deeper level. If you are acting to pursue your self-interest, can it also be that you are acting morally, or is it just a happy coincidence? Let's take another example.

Suppose you spend a great deal of time, effort, and energy helping a friend through a difficult time in her life simply because you miss going out with her and want to get back to how things used to be. Have you acted morally or not? This case relates to an important debate about **moral motivation**: Is an action "moral" only if it is done for moral reasons? We will return to this in a later chapter, for Immanuel Kant (1724–1804) claimed that an action has moral worth only if it stems from the appropriate moral motivation. But to stay focused on the main issue, recall that so far our central challenge is based on the claim that human beings are constituted to pursue their own self-interest, and nothing but their own self-interest. And we have conceded that morality often requires self-denial. From these premises it follows that human beings cannot always act morally. This is obviously a serious issue. So now let's directly evaluate the root of this argument: the psychological egoist claim that human beings must always act out of self-interest.

The Evidence for Psychological Egoism

Is it true that human beings are constituted to pursue their own self-interest and nothing else? Many intelligent and thoughtful people think so. And this is a good reason to take the claim seriously. But as we know, simply thinking something doesn't make it true. Neither does the fact that you (if you do), your parents, or some of your friends think it. None of us can make something true by thinking it is true (except in very special cases). We need reasons, in the form of evidence or arguments, to support claims to truth. What reason is there to believe that human beings are constituted to always follow their own self-interest? Psychological egoism is after all a psychological theory about human beings; it is something that should be supportable by evidence. It is a theory about human behavior. So what is the evidence that human beings are constituted to pursue nothing but their own interest?

Of course it is easy to find evidence of human beings acting in their own self-interest. If we look at the world of business, just reading the newspapers reveals breathtaking examples. But this isn't good enough. The theory is stated as a universal truth about all human beings and all of their actions. Logically, the claim "all human action is self-interested" is equivalent to the claim "no human actions are not self-interested." That is a bold suggestion indeed. The key issue, then, comes down not to how many positive instances can we find, but to finding proof that no negative instances exist: in this case, people engaging in genuinely self-sacrificing behavior. Compare the generalization that "all swans are white." We can come up with no end of examples of white swans, but nothing will prove the theory until we are somehow reassured that there are no swans of other colors (and famously, black swans were discovered in Australia eventually, thereby falsifying the previous theory).

Thus we need to turn the investigation on its head and look for examples of self-sacrificing, or self-denying, behavior among humans. Are there any? On the face of it, there are many. Maybe you've spent part of the weekend tidying up the apartment while

your roommates were out enjoying the sun. Is this not a form of self-denying behavior?

The psychological egoist is not going to be very impressed by this example. Maybe you detest an untidy apartment so much that you would rather tidy it up than go out and enjoy yourself. Maybe you are tidying today in the attempt to guilt-trip your roommates into cooking for you next week. In either case the psychological egoist can say that you are taking a short-term loss for longer-term gain. The example does not show that long-term self-denying behavior exists.

But really, do we want to say that pure self-denying behavior is impossible? Here is the German philosopher Arthur Schopenhauer (1788–1860), actually a defender of something close to psychological egoism, on the topic:

> [There] are individual but undoubted cases where not only punishment by law, but also discovery, and even any hint of suspicion were totally excluded, and yet a rich man had what belonged to him given back by a poor man: e.g., where something lost and found, or something deposited by a third party who had then died, was brought to its owner, or where something secretly left with a poor man by someone fleeing the country was faithfully kept and returned. (Schopenhauer, 1841/2009, p. 186)

True, Schopenhauer goes on to say that such people are as rare as "four-leaved clovers," but at this point all we need is a single example to refute the case.

Mother Teresa (1910–97), who was made a saint in September 2016, is often mentioned as an inspiring example to us all and a refutation of the self-interest theory. She spent much of her life tending to the sick and poor in Calcutta, at great hardship to herself. Was this all part of a calculated plan to advance her self-interest in the long term? What did she hope to gain? In fact some critics have been skeptical about Mother Teresa's achievements. But even if the commonly believed story is correct, the particular case is easy for the psychological egoist. There is, of course, a straightforward calculation of self-interest. Mother Teresa was a Christian. Christian doctrine includes belief in the afterlife. To ascend to heaven a person has to do good on earth, and so all apparently moral behavior is really a form of hidden self-interest. As we noted in the last chapter, there is an interesting theological question here: Is religious belief in an afterlife, or even God's reward on earth, a way of reducing morality to self-interest? If moral behavior is rewarded and immoral behavior punished, and people are motivated by the reward and punishment, then morality is no more than enlightened self-interest after all. If Mother Teresa's plan was to do good simply to get into heaven, then her behavior is consistent with the self-interest theory. Still, we have to wonder whether she had to do *so much* good for others. Couldn't she have relaxed a little?

But let's imagine an atheist version of Mother Teresa. Admittedly, the self-interest theorist could jump on this suggestion and reply that an imagined example is irrelevant. After all, even though we can imagine a green swan, it doesn't show that not all swans are white: We need to find an actual, rather than an imaginary, nonwhite swan to show that. For the moment let's continue with the thought experiment, accepting that on its own it doesn't prove anything. I introduce it not as a direct counterexample

to the self-interest theory, but as a stand-in for the many people who say they do not believe in the afterlife and yet also take a lot of trouble to do good for others. Let's use the idea of "Atheist Mother Teresa" to represent these people.

Once again, though, the self-interest theorist can make short work of Atheist Mother Teresa. True, Atheist Mother Teresa does a lot of good for others. But, so the theory goes, at some level she is acting this way because of the kick she gets out of it. It feels good to help, and those who help others do it to get a "warm glow" of righteousness. What looks like an act of self-sacrifice is an act of selfishness after all.

Well, is that right? How would Atheist Mother Teresa respond if you put this case to her? No doubt she would feel pretty hurt and insulted; she would vehemently deny that she has chosen a life of poverty, struggle, and sacrifice because, like a cocaine addict, she gets a rush from it. But you can press the point. Doesn't it feel good to do good? "Yes," she might say, "but that's not why I do it. If I didn't think it was the right thing to do, I wouldn't do it. And if I didn't think it was the right thing to do, I would not get a warm glow." In any case, the warm glow she feels barely registers on the scales compared to the pains suffered through a life of self-sacrifice. In other words, even if there is a warm glow, it is a small positive reward in comparison to the suffering she also experiences and not the main motivation of the action.

Here the critics might accuse Atheist Mother Teresa of deceiving herself. They could claim that no one would behave like this unless the warm glow more than compensated for the pain. But why should we believe this? Or, to use a simpler example, consider an atheist member of a resistance movement against a totalitarian regime who allows herself to be tortured to death to avoid betraying the movement. No warm glow, surely, would be enough to compensate for a prolonged and horribly painful death. Such cases seemingly refute the claim that humans are incapable of self-sacrificing behavior.

But the critic will be unbowed. These are fictional cases. In real-life cases, how do we actually know what is going on in someone's head? Perhaps the warm glow is ecstatic. Perhaps all atheist revolutionaries have a deathbed conversion and believe they are going to eternal bliss, or they believe the reputation of glory is enough to compensate for the pain. The self-interest theorist can dig in and insist: All action *must* be for the sake of self-interest. Notice, though, that once this move is made, the argument has changed significantly. Initially the critic made a psychological argument that drew on claimed empirical evidence: Anyone who appears to be self-sacrificing actually has, deep down, some selfish ulterior motive that we can discover. When faced with imagined counterexamples, the critic has changed to a theoretical argument: Self-sacrifice is simply not possible, and there *must* be an ulterior motive. But why must there be? Doesn't this argument assume that psychological egoism is true, rather than provide any grounds for believing it is true?

It is hard to see, though, how it could be possible to make a definitive argument for the self-interest theory, or indeed find a

definitive refutation. The power of argument is sometimes more limited than we would hope. The psychological egoist, if utterly convinced of the truth of the view, can insist that the warm glow must outweigh any pain or sacrifice, even if this now sounds rather **dogmatic**. We can now appreciate, though, that the psychological egoist is simply someone who holds one theory—and perhaps it is not as plausible as it first looked. Psychological egoism needs support from evidence or argument, and so it is not obviously true. But even if the theory is not obviously true, that does not make it false.

Can Psychological Egoism Be Rejected?

Nevertheless, psychological egoism can be a hard view to shift. For example, in *The Republic*, Plato sets himself the task of trying to show that acting justly is not mad or foolish, even if you can get away with behaving however you want.

In a sense, though, Plato puts himself in a difficult position in the argument. What does he have to show? He thinks he needs to demonstrate that it is in people's interest to act morally, even if we can profit from immorality without fear of being caught and punished. How would you show that acting morally is in your interest even when you can get away with breaking the rules? We will not go through the details of Plato's argument because it depends on many other elements of Plato's philosophical system as a whole. Here we need to pay attention to one important claim he makes—that a certain special type of inner harmony comes from acting morally. A modern version of this view might be that if you acted very badly, you simply wouldn't be able to live with yourself, perhaps through the nagging and life-wrecking guilt or shame of having harmed others for your own benefit. If you act immorally, it will eat away at you and make you miserable. Suppose, for example, that you had the Ring of Gyges and went into a poor old couple's house—people who had only ever been kind to you—and took their last food because they live nearer to you than anyone else, and this is the easiest way for you to feed yourself. You know you will never be caught. Would you feel okay about it? If not, why not?

You might feel, however, that Plato's argument—or at least the modern version of it—reduces morality to something like self-interest. We have just seen the argument that the only reason people act morally is that doing so gives them a kick of some sort, a warm glow. Here we have an argument that is something like the opposite of the warm glow argument; what we might call the "cold chill" argument. It says that people avoid immorality because they don't want to suffer the cold chill of the remorse and guilt of having done something wrong. But then this theory also appears to reduce morality to self-interest, simply in a different way.

How might we avoid collapsing morality into self-interest? Think about the nice old couple whose food you stole when you had the Ring of Gyges. How would you feel if you had acted that way? Most of us, I believe, would feel awful and probably prefer never to think about it. Why? Is it because you have been brought up to obey ordinary moral rules and find them impossible to shake, like someone who

automatically follows the rituals of the religion of their childhood even though they no longer believe? Or is it because the poor old couple will now find it hard to feed themselves in the coming days and will suffer real hardship as a result, unless someone else comes to help them? Knowing that other people can suffer, just as you would, seems to give you a reason to treat them well if you can. If that is your reason to feel awful when you break the rules, then the rules of morality seem to involve something beyond self-interest.

SELF-INTEREST AND EVOLUTION

Psychological egoists can try to find support from another quarter, evolutionary biology. The biologist J. B. S. Haldane (1892–1964) reportedly said, "I'd lay down my life for two brothers or eight cousins." This was a calculation based on what is now known as the **selfish gene** theory, as developed in detail by the contemporary evolutionary biologist Richard Dawkins (b. 1941). The idea is that old-fashioned self-interest theory is too crude. Human beings aren't selfish, our genes are. In this view, we might well sacrifice ourselves, but only if it is likely to lead to the greater preservation of our genes. Note, though, that the terminology can be confusing. Dawkins is not proposing that all human beings have a "gene for selfishness" (even if some people might have). Rather, he suggests that we are genetically programmed to try to preserve our genes. And we are programmed to do this whether we alone possess those genes or whether we have them in common with other people, in which case we will want to advance the interests of those other people.

Selfish Genes and Kin Altruism

The selfish gene hypothesis can seem powerful. When we think about the people in our world who are most self-sacrificing, mothers, caring for their children, are at the top of the list. A mother can be sure that her child contains half the genetic material that she has. So if, at some deep level, a mother's main motivation is to ensure the continuation of her genes, then she will do a good deal to help ensure that her children survive and reproduce. She will also care about her own brothers and sisters, who also have half her genes. Her nieces and nephews have a quarter of her genes, and her cousins have an eighth. This is why Haldane said he would lay down his life for two brothers or eight cousins; either way, the same amount of genetic material would be preserved. And indeed this type of self-sacrificing behavior is routinely observed in the animal kingdom, especially among what have been called the "social insects" such as ants and bees. In some examples sterile insects—incapable of reproduction—work tirelessly for the good of the group, all of whom are close relatives.

The selfish gene hypothesis is connected to the theory of **kin altruism** (although historically the idea of kin altruism predated the term *selfish gene*). Essentially, kin altruism supposes that your selfish genes lead you to act altruistically to those who are related to you. The more closely you are related to these people, the more altruistic you will be; conversely, you will be less

altruistic toward people who are more distantly related.

In one way, though, the theory of kin altruism seems crazy. I do, of course, help my relatives from time to time. But I don't think I have ever reasoned that "I'd better help my brother, at some cost to myself, to make sure that my genes survive." We would be quite surprised, and possibly a little worried, if we came across someone who thought things through like that. The selfish gene hypothesis and the theory of kin altruism have to be understood as some sort of indirect motivation. One possible idea is that your genes cause you to love those who are closely related to you, and your love of family leads you to help them. The love is real, and it explains your behavior; but in this theory, your selfish genes strongly influence who you love. And we can concede, I think, that the kin altruism theory has some plausibility. We do seem to care most about those most closely related to us, and it would be no surprise if researchers found a genetic explanation for that phenomenon. In fact, the real surprise would be finding out this behavior has nothing to do with evolution or genetics.

What, though, is the scope of the theory? Is it intended to explain all altruistic behavior? It doesn't seem to explain why parents of adopted children devote such care and attention to them. Or why some people care much more for their cousins than for their brothers or sisters, or why others care for their friends much more than anyone who is related to them. Let's consider a husband and wife, who, we hope, are no more closely related than any two strangers taken at random. Kin altruism can explain why they care for each other if they have children together: If your husband or wife flourishes, this is likely to be good for your children and hence the survival of your genes. But suppose they have no children and are past childbearing age. Kin altruism can no longer explain why the spouses might make sacrifices for each other, although it might say that in all these cases our emotions are somehow tricked or hijacked. That may be true; but overall, kin altruism seems unable to directly explain a whole range of altruistic behavior, even if it seems to get a certain amount right. Worse still, kin altruism doesn't seem compatible with the behavior of family members who sometimes harm or neglect each other. Some will think that kin altruism gets things backward. What motivates us to help family members, they will claim, is not our genes but the love and close bonds we develop in growing up with our family. And the development of mutual affection will also explain why we make sacrifices for family members who are not genetically related.

The Mountain People

Besides the selfish gene theory, the evolutionary theorist has another card to play. We will begin discussing this next idea by looking at what was claimed to be a real-life example—although as we will see, we should be highly skeptical. In 1972 the anthropologist Colin Turnbull (1924–94) published a sensational book, *The Mountain People*, about a group called the Ik people who live in the mountains between Uganda, the Sudan, and Kenya. According to Turnbull, these people had virtually no

compassion or feeling for each other. Some of his examples sound like cartoon horror stories, giving accounts of mothers laughing when their children burn themselves. Here is one of the most dramatic:

> The mother goes about her business [at the water hole or in the fields] leaving the child there, almost hoping that some predator will come and carry it off. This happened once when I was there—once that I know of, anyway—and the mother was delighted. She was rid of the child and no longer had to ... feed it, and still further this meant that a leopard was in the vicinity and would be sleeping the child off and thus be an easy kill. (Turnbull, 1972, p. 136)

The author gives us many other examples. Old people were neglected, abused, even killed. Food is taken from the mouths of the starving. The situation of the tribe seems unbearably bleak. But on reading the book, a question strikes you again and again. If this is how they behaved to each other, how did the society survive? Turnbull's answer is that this behavior is relatively new to the tribe. Since being forced off their traditional hunting grounds, which were turned into a highly protected nature reserve, the Ik people were now starving to death although they pretended not to be. Their egoistic behavior was an aberration; if it lasted, the tribe would not survive. Anthropologists now completely reject the credibility of Turnbull's account, arguing that his study was deeply flawed. Indeed it is said that upon hearing about the way Turnbull had represented them, the Ik people considered taking legal action against him, thus showing a level of sophistication that does not seem to fit with Turnbull's account. Rather than as a depiction of fact, then, let's treat Turnbull's book as another example of a "thought experiment." Could a tribe like the Ik ever have existed? Could they reproduce over the generations?

The answer is surely that such a tribe could not have survived. As Turnbull himself noted, he expected that the Ik tribe would soon die out. Group survival needs a measure of altruism, of individual self-sacrifice. Hence another evolutionary argument seems possible: Groups that don't develop cultures of altruism don't survive. And so a mechanism known as group selection has been suggested. This is Charles Darwin (1809–82) in his book, *The Descent of Man*, on the matter:

> A tribe including many members who, from possessing in a high degree the spirit of patriotism, fidelity, obedience, courage, and sympathy, were always ready to aid one another, and to sacrifice themselves for the common good, would be victorious over most other tribes; and this would be natural selection. (Darwin, 1871/2004, p. 157)

How, though, is this meant to work? Evolutionary biologists distinguish between three different possible moral evolutionary mechanisms. We have already discussed kin altruism (the theory of the selfish gene). The Ik and Darwin lead us to another possibility —**group altruism**: Groups die out if they don't develop altruism. And finally we need to look at the idea of **reciprocal altruism**: "If you'll scratch my back, I'll scratch yours." This is a fairly local, small-scale mechanism in which people implicitly agree to help each other. These three ideas—kin, group, and reciprocal altruism—seem to be closely related: Groups are likely to be made up of people who are related and who do mutual favors for each other. Group altruism suggests that individuals will engage in

behavior that favors the group, even at a cost to themselves. Reciprocal altruism is rather more fine-grained; it occurs when individuals within the group cooperate with other particular individuals, but only when they also cooperate back from time to time.

Some evolutionary biologists, however, have questioned whether there is such a thing as pure group altruism, unless it is backed by kin or reciprocal altruism. It is plausible that groups with moral codes will survive and outcompete with groups without a moral code holding them together, but the problem is that a selfish and clever individual living with a group of altruists will have a tremendous advantage. He or she will be able to exploit others, and if this advantage turns into more offspring, the gene for selfishness (remember, not the same thing as the selfish gene) will be passed on and eventually dominate the group. Accordingly, some have argued that altruism within a group is not stable unless backed up with either kin altruism or reciprocal altruism (or both) to put a brake on **free riding**. Others have argued that in certain ways, group altruism can resist some level of noncooperative selfish behavior.

Reciprocal altruism, nevertheless, seems likely to be a stronger mechanism. It is the idea that I will look after you in the expectation that you will look after me. This can explain why nongenetically related individuals such as a husband and wife, or a group of friends, will help each other. Once more, this theory must contain a good deal of truth. People with a reputation for being selfish tend to get left out of group activities and are worse off for that. Therefore, we all have self-interested reasons for cooperating with at least the people we see regularly. On the other hand, it seems that what matters in the first instance is the reputation, not the reality. If you can be a free rider**,** pretending to be cooperative but slyly getting away with what you can when you can—as if you had the Ring of Gyges—then that's what self- interest says you should do. But in a small group, or a group that regularly interacts, free-riding often is detected and punished. (This may explain why villages are generally safer and more secure than large cities: Uncooperative behavior is much more easily noticed and dealt with.) In Chapter 7 we will return to the issues of free-riding.

To conclude the discussion of psychological egoism, even if it is true that morality can be reduced to self-interest, the task of the theorist of morality is not complete. The reason for this is that moral theory comes in many forms and varieties, and it is not clear which form of morality would best promote self-interest. Moral philosophy is needed even if in a somewhat different form than we normally suppose. It leads to the question of the optimum mix of short-term self-interest and altruism to preserve long-term self-interest.

ETHICAL EGOISM

So far we have been looking at varieties of what we called psychological egoism, which as I noted, is essentially, a scientific theory about human behavior. We have seen some arguments for the conclusion that human beings are somehow constituted or genetically programmed to pursue their own

self-interest; and this, at least at first, raises the question of whether morality is possible. But we have also seen that once evolutionary arguments are taken into account, in some ways psychological egoism can generate behavior that looks very similar to ordinary morality.

But we need to explore another view, which I referred to above as ethical egoism. It says that whether or not psychological egoism is true, following your own self-interest is the *morally right thing* for you to do. Ethical egoism comes in many different forms, for "self-interest" can be understood in many different ways. Normally we think of self-interest in a relatively narrow form—getting what you most want, satisfying your desires, or perhaps achieving financial success and having a life of luxury. These are what we could call your "selfish" self-interests. But we can also take a broader view of self-interest by arguing that, for example, doing good things for others is part of an individual's self-interest if that is what she most wants to do. In this view a type of harmony exists between the person's self-interest and the good of others; and so by acting apparently selfishly, she is automatically acting morally. … Aristotle held a view of this sort, and we will discuss his theory in detail. For now, however, let's restrict ourselves to what we could call "narrow ethical egoism," which claims that the morally right thing to do is to pursue your selfish interest. But can this really be right? Doesn't morality call on us to make, say, a small sacrifice to save a life? How could it be thought otherwise?

Private Vices, Public Virtues

Narrow ethical egoism comes in at least two versions. Consider this famous quote from the Scottish economist and moral philosopher Adam Smith (1723–90) from his book *The Wealth of Nations*:

> It is not from the benevolence of the butcher, the brewer, or the baker that we expect our dinner, but from their regard to their own interest. We address ourselves, not to their humanity but to their self-love, and never talk to them of our own necessities but of their advantages. (Smith, 1776/1982, p. 119)

Smith points out a remarkable fact about exchange in market economies. Economic agents are driven by self-interest: Both sides want the best deal they can get. But for butchers, as an example, it is in their self-interest to give customers good value so that they come back another time. In the longer term the butchers can get what they want—a regular profit—only by giving customers what they want—good meat. To pursue their self-interest, butchers have to give customers excellent service. In this way the pursuit of individual self-interest leads to the collective common good. As the Anglo-Dutch thinker Bernard Mandeville (1670–1733) pointed out, "private vices" lead to "public virtues" (Mandeville, 1714/1989).

This form of ethical egoism is, in a sense, conditional. It tells us to follow our own self-interest, but only because in this way, everyone's interest in the common good will be advanced. Some defenders of the capitalist free market economy have been impressed by this argument. Opponents of the free market, such as socialists, criticize capitalism because it encourages selfishness.

But in reply, defenders say that capturing selfishness for the public good is actually the strength of capitalism. You can make a profit only by giving people what they want, so the greedier the capitalist, the better it is for the customer. Prices come down and quality goes up in attempt to capture more of the market. On this basis it is often argued that the way to advance the collective good is to allow the capitalist economy to function with minimal government intervention or regulation. This, it is said, will allow individual selfishness to flow for public benefit.

Unfortunately, as Adam Smith was himself aware, his argument has much more limited scope than some of his followers have suggested. Consider, for example, a financial agent selling you a pension plan. It is in the narrow self-interest of the agent to convince you that the right plan for you is the one that gives him the greatest commission. But if so much commission goes to the agent, it is rather likely that the product he is boosting will give you a poor deal. The same is true for agents selling subprime mortgages, product insurance, collateralized debt, or any other sophisticated financial product. Their interest is to lead you to buy whatever gives them the highest profit. If the product is no good, then, unlike the meat you bought from the butcher and eat that evening, it could be years before you find out. By that time the agent, most likely, will have moved on and be trying to sell something else. And this is no mere theoretical issue. In the view of some analysts, the root cause of the financial crisis of 2008 was a belief by some financial regulators that Adam Smith had shown that the pursuit of self-interest in an unregulated market would always work out in the common interest. But as we learned at enormous cost, Smith himself knew he had shown no such thing.

Pure Ethical Egoism

To recap, I said that ethical egoism comes in two forms. In the form we have been looking at so far, it states that pursuing your own self-interest is the best way of advancing the common good. I have given readers some reason to believe that, although this theory has merit in limited contexts, as a generalization it is highly problematic. But in any case, when understood this way, ethical egoism seems to fall short of being an "ultimate" moral theory: It presupposes that we can reach an independent "common good" by acting in our self-interest. It is a theory about how to achieve good results, understood in non-egoistic terms. Therefore, strictly speaking, it is not an egoist theory at all.

The second form of ethical egoism is uncompromising. Rather than arguing that we have a duty to act in a self-interested way because of the consequences, it suggests that it is right to act in our own interests, *whatever* the consequences for other people. This theory, sometimes associated with the influential Russian American thinker and novelist Ayn Rand (1905–82), has been described as arguing for the "duty of selfishness." This slogan reflects a striking and bold theory, apparently suggesting that it is your moral duty to be as grasping and selfish as you can be. But at least in some of her writings,

Rand makes clear that her principal argument is against a type of joyless, altruistic self-sacrifice that requires individuals to put their own interests aside to promote the well-being of strangers. She clearly indicates that she believes a good life—a selfish life—can include generous behavior toward your family and friends, the enjoyment of which also contributes to your own self-interest. In this respect, her view allows a level of harmony between self-interest and action for the sake of others, rather than instructing everyone to follow a narrow, calculating path. Yet there remains a powerful element of self-concern in this view, given the strong stand Rand takes regarding some forms of self-sacrifice. Conventional morality, thinks Rand, is our "enemy," for we can only lose from it. In *The Virtue of Selfishness*, she says:

> By elevating the issue of helping others into the central and primary issue of ethics, altruism has destroyed the concept of any authentic benevolence or good will among men. It has indoctrinated men with the idea that to value another human being is an act of selflessness, thus implying that a man can have no personal interest in others—that *to value* another means *to sacrifice* oneself—that any love, respect or admiration a man may feel for others is not, and cannot be a source of his own enjoyment, but is a threat to his existence, a sacrificial blank check signed over to his loved ones. (Rand, 1964, pp. 49–50)

These are strong words. Yet we need to be careful. Rand is here rejecting a view of morality that claims moral value is found only in promoting the interests of others at your own cost. Certainly, some conceptions of morality do require such a level of sacrifice, such as extreme religious views. Or consider Mrs. Jellyby, a character created by Charles Dick- ens (1812–70) in his novel *Bleak House* (1853), to show how unappealing extreme altruism can be. Mrs. Jellyby lets her children go hungry and dirty while she devotes her time to setting up a mission in Africa. But rather than abject self-sacrifice, most moral views seek some sort of balance between the interests of the individual, their family, and other people. So rejecting such an extreme form of altruism is compatible with endorsing a wide range of approaches to morality, and not just Rand's own view. It would be a logical fallacy to think that rejecting extreme altruism results in establishing a form of ethical egoism, because intermediate positions do remain possible.

Nevertheless, the idea that pursuing your own selfish interests is the morally right thing to do has been discussed since the beginnings of moral philosophy. Ethical egoism, in this pure form, may look like the opposite of a moral theory because it seemingly instructs you to be indifferent to the plight of other people, especially strangers, however bad their situation is—or, indeed, whatever damage you do to them. Pushed to the extreme, this behavior would permit or perhaps require you to rob or even kill people in pursuit of your interest.

Ethical egoism is thus a highly counterintuitive theory. But its appeal is easy to see, at least in more moderate cases. It starts, most naturally, from facts about human existence. You have your own life, your own perspective on the world. And you will have only one life. No one else can feel your pleasures or pains, and you cannot feel those of anyone else. Although others may

occasionally help you, or appear to do so, in the end the only person you can rely on is yourself. You should, then, do what you can to make your life as fulfilling and worthwhile as possible. This is an obligation to yourself; a moral duty to advance your own interests. Of course you can choose to help people, but you must not neglect yourself. You were given this one life to lead, and it is your duty to make the best of it.

Put like this, ethical egoism is not merely a type of convenient moral screen for selfishness. Instead, it seems to reflect a deep fact about human existence: Each of us has our own life to lead, and we are responsible for making our own life go well. This theory seems to respect individual responsibility, and in some ways seems to be of great rigor and integrity. However, its difficulties are also easy to see: First, although it is true that I have a life to lead, so do you; and so does everyone else. Why should my life take priority over yours? Rather than potentially crushing others in pursuit of our own interests, it seems we need to respect everyone's right to pursue a meaningful and fulfilling life. From this point of view, ethical egoism is a very one-sided position: It addresses only half of the moral problem. The moral dilemma we often face is how to reconcile our own interests with those of others. We ignore half the dilemma if we simply assume that we have a moral duty to pursue our own interests. And this point brings a second problem into focus: Even if we do have a moral duty to pursue our own interest, we can see how badly that can go by remembering the horror stories of the Ik. Perhaps, then, pursuing our own interests requires significant compromise, to avoid what Thomas Hobbes (1588–1679), in his book *Leviathan*, called "a warre ... of every man against every man" (Hobbes, 1651/1997, p. 61). …