Objections to Utilitarianism and Responses

MacAskill, William, et al.

> *Bernard Williams (...) concluded a lengthy attack on utilitarianism by remarking: 'The day cannot be too far off in which we hear no more of it.' It is now more than forty years since Williams made that comment, but we continue to hear plenty about utilitarianism.*
> - Peter Singer & Katarzyna de Lazari-Radek

Utilitarianism is a very controversial moral theory. Critics have raised many objections against it, and its defenders have responded with attempts to defuse these objections.

In the following, we explain and address the most influential objections to utilitarianism.

General Ways of Responding to Objections to Utilitarianism

Many objections rest on the idea that utilitarianism has counterintuitive implications. We can see these implications by considering concrete examples or *thought experiments*. For instance, in our article on the rights objection against utilitarianism, we consider the Transplant case:

> Imagine a hypothetical scenario in which there are five patients, each of whom will soon die unless they receive an appropriate transplanted organ—a heart, two kidneys, a liver, and lungs. A healthy patient, Chuck, comes into the hospital for a routine check-up and the doctor finds that Chuck is a perfect match as a donor for all five patients. Should the doctor kill Chuck and use his organs to save the five others?

At first glance, it seems that utilitarianism has to answer the question with "Yes, the doctor should kill Chuck." It is better that five people survive than that just one person does. But on commonsense morality and virtually every other moral theory, the answer is "No, do not kill Chuck." On most views, killing Chuck would be morally monstrous. This apparent counterintuitive implication of utilitarianism is taken as an argument against its being the correct moral theory.

Proponents of utilitarianism can respond to the counterintuitive implications of utilitarianism in cases like Transplant in four separate ways.

First, they can *accommodate* the intuition that seems to conflict with the utilitarian recommendation by emphasizing that a sophisticated, as opposed to naive, application of utilitarian principles avoids the counterintuitive implication. We usually lack detailed information about the future consequences of our actions. This suggests that usually it would be naive for us to act in a way deemed very wrong on commonsense morality, under the impression this will bring about the best outcome. A sophisticated utilitarian—in line with multi-level utilitarianism—recognizes the limitations to our foresight and acts in accordance with commonsense heuristics, other than in exceptional circumstances.

Second, utilitarians can attempt to *debunk the moral intuition* invoked by a particular case by pointing out that it resulted from an unreliable process.[2] If a debunking argument succeeds, the respective moral intuition should not be given much weight in our moral reasoning. As an illustration, consider that in most Western societies Christianity was the dominant religion for over one thousand years, which explains why moral intuitions grounded in Christian morality are still widespread. For instance, many devout Christians have strong moral intuitions about sexual intercourse, which non-Christians do not typically share, such as the intuition that it is wrong to have sex before marriage or that is wrong for two men to have sex. The discourse among academics in moral philosophy generally disregards such religiously-contingent moral intuitions. Many philosophers, including most utilitarians, would therefore not give much weight to the Christian's intuitions about sexual intercourse.

Third, proponents of utilitarianism can *attack the available alternatives*—such as deontological or virtue ethical theories—to show that they, too, have implications no less counterintuitive than those of utilitarianism.

A fourth strategy is to *tolerate* the intuition, which is sometimes called "biting the bullet". This is to accept that utilitarianism has counterintuitive implications but then holding on to the theory because all-things-considered it is still more plausible than its rivals. The costs of accepting a counterintuitive implication, it is argued, can be outweighed by the benefits of the arguments in favor of utilitarianism. Moreover, it is impossible to have any non-ad-hoc theory that accords with all of our intuitive moral judgements. Our intuitions are often inconsistent and they are subject to change over time, which makes it impossible to find consistent and plausible principles that reflect all of them. So it requires judgment to determine which intuitions and theoretical commitments are non-negotiable, and which we should be willing to compromise on in pursuit of "reflective equilibrium", or the most plausible and coherent overall combination of moral judgments and principles.

Specific Objections Against Utilitarianism

In separate articles, we discuss the following critiques of utilitarianism:

(1) The rights objection charges utilitarianism with being overly *permissive*. To maximize the sum total of well-being, utilitarianism might be thought to allow infringing upon others' rights or violating other apparent moral constraints.

(2) The demandingness objection claims that utilitarianism is overly *demanding* because it requires excessive self-sacrifice from us in order to help others.

(3) The equality objection holds that utilitarianism is not sufficiently concerned with distributive justice and the value of *equality*.

(4) The mere means objection claims that utilitarianism treats people as "mere means" to the greater good rather than as ends in themselves.

1. The Rights Objection

According to commonsense morality and many non-utilitarian theories, there are certain *moral constraints* you should never violate. These constraints are expressed in moral rules like "do not lie!" and "do not kill!". These rules are intuitively very plausible. This presents a problem for utilitarianism. The reason for this is that utilitarianism not only specifies which outcomes are best—those having the highest overall level of well-being—but also says that it would be wrong to fail to realize these outcomes.

Sometimes, realizing the best outcome requires violating moral constraints against harming others—that is, violating their rights. There is no reason to expect commonsense moral rules to always coincide with the best ways to act according to utilitarianism; sometimes they conflict. An example of such a conflict is the Transplant thought experiment:[1]

> **Transplant:** Imagine a hypothetical scenario in which there are five patients, each of whom will soon die unless they receive an appropriate transplanted organ—a heart, two kidneys, a liver, and lungs. A healthy patient, Chuck, comes into the hospital for a routine check-up and the doctor finds that Chuck is a perfect match as a donor for all five patients. Should the doctor kill Chuck and use his organs to save the five others?

At first glance, it seems that utilitarianism has to answer the question with "Yes, the doctor should kill Chuck". It is better that five people live than that one person lives. But on commonsense morality and virtually every other moral theory, the answer is "No, do not kill Chuck". On most views, killing Chuck would be morally monstrous. Utilitarianism seems to be the rare exception that claims otherwise. This apparent implication is often taken as an argument against utilitarianism being the correct moral theory.

Proponents of utilitarianism might respond to this objection in four ways. We will go through them in turn.

Accommodating the Intuition

A first utilitarian response to the thought experiment might be to *accommodate the intuition* against killing Chuck by showing that utilitarianism does not *actually* imply that the doctor should kill him for his organs. Critics of utilitarianism assume that, in Transplant, the doctor killing Chuck will cause better consequences. But this assumption may be questioned. If the hospital authorities and the general public learned about this incident, a major scandal would result. People would be terrified to go to the doctor. As a consequence, many more people could die, or suffer serious health problems, due to not being diagnosed or treated by their doctors. Since killing Chuck does not clearly result in the best outcome, and may even result in a terrible outcome, utilitarianism does not necessarily imply that the doctor should kill him.

Even if we stipulate that the scenario is an unusual situation in which killing Chuck *really would* lead to the best outcome (with no further unintended consequences), it is hard to imagine how the *doctor* could be so certain of this. Given how incredibly bad it would be to undermine public trust in our medical institutions (not to mention the reputation harm of undermining utilitarian ethics in the broader society),[2] it would seem unacceptably *reckless*, according to expectational utilitarianism, for the doctor to risk such population-wide harm to save just a small handful of lives. Utilitarianism can certainly condemn such recklessness, even while allowing that there are rare cases in which, by unpredictable fluke, such reckless behavior could turn out to be for the best.

This is a generalizable defense of utilitarianism against a wide range of alleged counterexamples. Such "counterexamples" invite us to imagine that a typically-disastrous class of action (such as killing an innocent person) just so happens, in this special case, to produce the best outcome. But the agent in the imagined case generally has no good basis for discounting the typical risk of disaster. So it would be unacceptably risky for them to perform the typically-disastrous act.[3] We maximize expected value by avoiding such risks.[4] For all practical purposes, utilitarianism recommends that we should refrain from rights-violating behaviors.

Debunking the Intuition

A second strategy to deal with the Transplant case is to debunk the intuition against killing Chuck by showing that the intuition is unreliable. A utilitarian might argue that it is almost always wrong to commit murder and that we should cultivate strong character dispositions and social norms against murder. Therefore, our intuition against killing Chuck may just result from us having embraced a general moral norm against murder. While this norm is correct in the vast majority of cases, it can fail under those very exceptional circumstances where killing someone would actually bring about the best consequences.

We may also worry that the intuition reflects an objectionable form of status quo bias. However terrible it is for Chuck to die prematurely, is it not—upon reflection—equally terrible for any *one* of the five potential beneficiaries to die prematurely? Why do we find it so much easier to ignore their interests in this situation, and what could possibly justify such neglect? There are practical reasons why instituting rights *against being killed* may typically do more good than rights *to have one's life be saved*, and the utilitarian's recommended "public code" of morality may reflect this. But when we consider a specific case, there's no obvious reason why the one right should be more important (let alone five times more important) than the other, as a matter of principle. So attending more to the moral claims of the five who will otherwise die may serve to weaken our initial intuition that what matters most is just that Chuck not be killed.

Attacking the Alternatives

A third response to the Transplant case is to *attack the available alternatives* to utilitarianism to show that they have even more counterintuitive implications.

All of the standard arguments against deontic constraints become relevant at this point. For example, the hope objection flags that a benevolent observer should prefer that the five be saved, and it's hard to see how deontic moral rules could matter more (or have greater normative authority) than what we—or any impartial benevolent observer—should hope is done.

As noted above, the charge of status quo bias seems especially pressing in this context. If you asked all six people from behind the veil of ignorance whether you should kill one of them to save the other five, they'd all agree that you should. A 5/6 chance of survival is far better than 1/6, after all. And it's morally arbitrary that the one happens to have healthy organs while the other five do not. There's no moral reason to privilege this antecedent state of affairs, just because it's the status quo. Yet that's just *what it is* to grant the one a right not to be killed while refusing the five any rights to be saved. It is to arbitrarily uphold the status quo distribution of health and well-being as morally privileged, no matter that we could improve upon it (as established by the impartial mechanism of the veil of ignorance).

Another challenge may be presented by increasing the stakes in our thought experiment:

> **Revised Transplant:** Suppose that scientists can grow human organs in the lab, but only by performing an invasive procedure that kills the original donor. This procedure can create up to one million new organs. Like before, our doctor can kill Chuck, but this time use his body to save one million people. Should she do this?

Consider how two non-utilitarians would react to Revised Transplant. The *Moderate non-utilitarian* says that, unlike in the original case, the doctor should kill Chuck because the constraint against harming others is outweighed, since enough is at stake. The *Absolutist non-utilitarian*, on the other hand, says that the doctor still should not kill Chuck, since no amount of benefit can outweigh the injustice of killing him.

One objection to the Moderate is that their position is *incoherent*. The rationale underlying the intuition that the doctor refrain from killing Chuck in Transplant should also forbid killing him in Revised Transplant. In both cases, an innocent person is sacrificed for the greater good. Another objection to the Moderate is that their position is *arbitrary*. The Moderate must draw a line past which constraint violations become permissible, for example, when the benefit is for at least one million people. But why draw the line precisely at that point, rather than higher or lower? What is so special about this particular number, 1,000,000? Yet the same question can be asked for any specific number of lives saved. The only non-arbitrary positions are that of the Absolutist, for whom there is no number of lives saved that can justify killing Chuck, and that of the utilitarian, who says that killing Chuck is justified if the benefits outweigh the costs.

The problem with Absolutism is this position is even more counterintuitive than utilitarianism. If we continue to increase the number of lives we could save by killing Chuck—say, from one million to one billion, and so on—it soon becomes absurd to claim that doing so is impermissible. This position appears even more absurd when we consider cases involving uncertainty. For instance, it seems the Absolutist is committed to saying it is impermissible to perform the medical procedure on Chuck, even if it had only a very small chance of killing him and is guaranteed to save millions of lives.

Tolerating the Intuition

The final response is for the advocate of utilitarianism to “bite the bullet”, holding on to the claim we should—in this hypothetical situation—kill Chuck despite the intuition that killing Chuck is wrong. It is regrettable that the only way to save the five other people involves Chuck’s death. Yet the right action may be to kill him since it allows the five others to continue living, each having meaningful experiences and enjoying their lives as much as Chuck would have enjoyed his own. Chuck’s death, while unfortunate, is *stipulated* by the thought experiment to be required to create a world where there is as much well-being as possible.

Of course, it's important to stress that real life comes with no such stipulations, so in real-life cases utilitarians overwhelmingly opt to "accommodate the intuition" and reject the assumption that killing innocent people leads to better outcomes.

2. The Demandingness Objection

Many critics argue that utilitarianism is too demanding, because it requires us to always act such as to bring about the best outcome. The theory leaves no room for actions that are permissible yet do

not bring about the best consequences. Philosophers label such actions that are morally good but not required as supererogatory; maximizing utilitarianism (like maximizing consequentialism more broadly) denies that any action can be supererogatory. As a result, some critics claim that utilitarianism is a morality only for saints.[1]

Very few people, including utilitarian philosophers, live their life in perfect accordance with utilitarianism. For instance, consider that the money a person spends on dining out could pay for several bednets, each protecting two children in a low-income country from malaria for about two years.[2] From a utilitarian perspective, the benefit to the person from dining out is much smaller than the benefit to the children from not having malaria, so it would seem the person has acted wrongly in choosing to have a meal out. Analogous reasoning applies to how we use our time: the hours someone spends on social media should apparently be spent volunteering for a charity, or working harder at one's job to earn more money to donate.

To many people, these extreme obligations of utilitarianism seem absurd at first glance. According to commonsense morality, we are permitted to spend most of our income on ourselves, our loved ones, and on our personal projects. Charity, by commonsense lights, is good and praiseworthy, but not obligatory.

Proponents of utilitarianism might respond to this objection in four ways. We will go through them in turn.

Accommodating the Intuition

One way to soften the demands from utilitarianism is to argue that morality should consider human psychological limitations, such as our weakness of will. Utilitarianism recognizes that we cannot work all the time to help others without burning out, which would lead to us doing less overall good in the long run. Similarly, we need to spend money on ourselves to stay reasonably happy and healthy to sustain our long-term motivation to do good.

In addition, it is often justified for utilitarians to spend money or time to accommodate the expectations and needs of other people. If utilitarianism is associated with extreme self-sacrifice, others may not want to join utilitarian causes. Likewise, it may sometimes be justified on utilitarian grounds to buy expensive dinners if that allows one to have valuable meetings with non-utilitarians who would not want to self-sacrifice.

However, even if we accept that spending resources on ourselves is of great instrumental importance for us to be able to benefit others, most of us must admit that we could be doing more. Utilitarianism remains a demanding ethical theory in practice, even when we account for the psychology of ourselves and others.

A more robust accommodation may be secured by rejecting the ordinary notion of moral "requirement". As explained in an earlier chapter:

> Utilitarians agree that you *ideally* ought to choose whatever action would best promote overall well-being. That's what you have the *most* moral reason to do. But they do not recommend *blaming* you every time you fall short of this ideal. As a result, many utilitarians consider it misleading to take their claims about what *ideally ought* to be done as providing an account of moral "rightness" or "obligation" in the ordinary sense.

According to utilitarianism, whether someone should be blamed for their actions is itself something to be decided by the consequences that blaming them would have. Blaming people whenever they fail to do the most good will likely have bad consequences, because it discourages people from even trying. Instead, utilitarianism will generally recommend praising people who take steps in the right direction, even if they fall short of the utilitarian ideal. This shows how the utilitarian notion of "wrongness" comes apart from the commonsense understanding of "wrongness", which is much more tied to blameworthiness.

Indeed, on a scalar or satisficing version of utilitarianism, doing less than the best need not be considered "wrong" at all. It is simply less than would be ideal. Satisficing utilitarianism identifies some lower minimum threshold for what is "required" to avoid blameworthiness, whereas scalar consequentialism eschews such thresholds entirely, instead assessing the moral quality of actions on a continuous scale from better to worse. It is better for an affluent person to donate 10% of their income to charity than to donate only 1%, which itself is better than donating nothing at all.

Debunking the Intuition

The second line of response is to argue that ordinary demandingness intuitions *presuppose* (rather than independently *support*) non-consequentialism. By asking the comparatively wealthy to do a lot to help the less fortunate, utilitarianism imposes some non-trivial costs on the wealthy. But compare this to the harms endured by the less fortunate by the wealthy doing less (or nothing) to help them. These harms vastly exceed the costs that utilitarianism would impose on the wealthy. Utilitarians may thus argue that it is the *non-utilitarian* views that are "too demanding" since they impose greater overall costs and focus these costs on those who are least able to bear them.

David Sobel develops this argument in The Impotence of the Demandingness Objection:[3]

> Consider the case of Joe and Sally. Joe has two healthy kidneys and can live a decent but reduced life with only one. Sally needs one of Joe's kidneys to live. Even though the transfer would result in a situation that is better overall, the Demandingness Objection's thought is that it is asking so much of Joe to give up a kidney that he is morally permitted to not give. The size of the cost to Joe makes the purported moral demand that Joe give the kidney unreasonable, or at least not genuinely morally obligatory on Joe. Consequentialism, our intuitions tell us, is too demanding on Joe when it requires that he sacrifice a kidney to Sally.
>
> But consider things from Sally's point of view. Suppose she were to complain about the size of the cost that a non-Consequentialist moral theory permits to befall her. Suppose she were to say that such a moral theory, in permitting others to allow her to die when they could aid her, is excessively demanding on her. Clearly Sally has not yet fully understood how philosophers typically intend the Demandingness Objection. What has she failed to get about the Objection? Why is Consequentialism too demanding on the person who would suffer

> significant costs if he was to aid others as Consequentialism requires, but non-Consequentialist morality is not similarly too demanding on Sally, the person who would suffer more significant costs if she were not aided as the alternative to Consequentialism permits?[4]

We may cast further doubt on our demandingness intuitions by noting other apparent inconsistencies in their application. For instance, many philosophers—utilitarian and non-utilitarian alike—would readily accept that morality can be very demanding in wartime. Under the circumstances of war they might think that people may have to make great sacrifices, including giving up their property or even their life. Yet in peacetime today hundreds of millions of people live in dire circumstances of extreme poverty, and billions of animals suffer in factory farms and are killed every year. At the same time, many affluent people enjoy a wide range of luxury goods and have access to effective channels through which they could assist the poor. From the utilitarian perspective, the world today is just as high-stakes as it is in wartime. For this reason it is no more demanding—and arguably much less—to require the affluent to donate money to assist the poor in the present day than it is to require soldiers to sacrifice their lives in a war against, say, cruel authoritarianism.

It is worth noting that utilitarianism would not ask as much of us if most affluent individuals acted morally and shared more of their resources with those most in need. Utilitarianism only becomes so demanding because few affluent people do anything significant to address the major problems in the world. However, even if the affluent were so charitable that there was negligible value in further charity, there would still be occasions when utilitarianism would make great demands of one, such as when one must sacrifice one's own life to save the lives of several others. Therefore, while the present point mitigates the force of the demandingness objection somewhat, it does not disarm it completely.

Attacking the Alternatives

A third response is to argue that non-utilitarian moral views are often *insufficiently demanding*. We have already established that citizens of affluent countries can prevent a substantial amount of suffering and death in developing nations at a comparably low cost to themselves by donating to highly effective aid organizations. According to many non-utilitarian views, it is good but entirely *optional* to donate a significant portion of our income to charity. However, this is arguably not demanding enough since it entails that we are not required to save lives even when we can do so at a low cost to ourselves. Therefore, these views violate Peter Singer's intuitively plausible assertion that "if it is in our power to prevent something bad from happening, without thereby sacrificing anything of comparable moral importance, we ought, morally, to do it".[5] As Singer explains, most people agree that it would be morally monstrous to just watch a child drown in a shallow pond when you could easily save them at the cost of ruining your expensive clothes. Saving innocent lives is worth some moderate financial cost, and any reasonable moral theory needs to reflect that fact.

In addition, proponents of utilitarianism can note that non-utilitarian views are sometimes even more demanding. Recall Sobel's example involving Joe and Sally. Commonsense ethics prohibits Sally from stealing one of Joe's kidneys, even if that would be the only way to save her own life (and the harm to Joe would only be moderate). This shows that commonsense morality can be very demanding sometimes, even requiring you to give up your life on moral grounds. While utilitarianism makes *different* demands from other moral theories, the demands of utilitarianism are not obviously less reasonable. They always have a good principled basis, after all.

Tolerating the Intuition

Finally, proponents of utilitarianism may once again "bite the bullet" and simply accept that morality is very demanding. They may point out that utilitarian demands are grounded in the compelling goal to create a flourishing world with as much well-being as possible for everyone. Whenever utilitarianism requires us to give up something we value to benefit others, at least we know that this benefit is greater, often much greater, than the cost to us.

3. The Equality Objection

Some argue that utilitarianism conflicts with the ideal of equality.

Suppose, for example, that you could choose between two possible distributions of well-being, *Equality* and *Inequality*: Equality has 1,000 people at well-being level 45, while Inequality has 500 people at 80 well-being and another 500 people at 20 well-being.

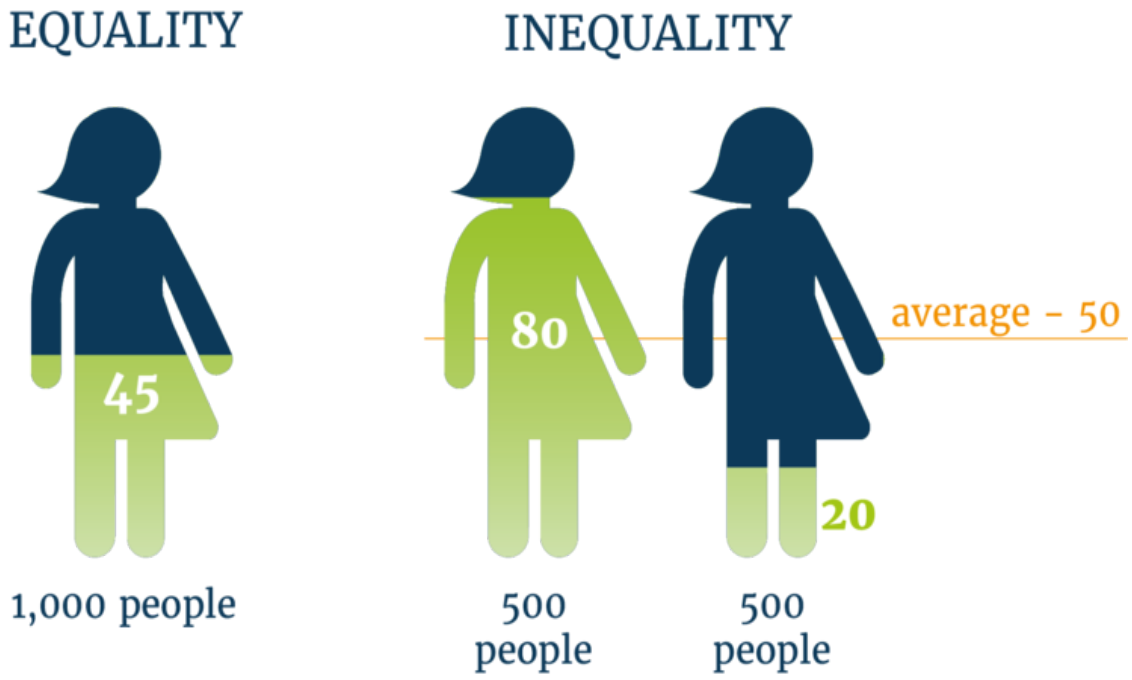

By the lights of utilitarianism, only the sum total of well-being determines the goodness of an outcome: it does not matter how that well-being is distributed across people. Since the sum total of well-being is greater in Inequality than in Equality, the unequal outcome is preferable according to utilitarianism.

Some philosophers object to the utilitarian view regarding this choice, claiming that the equal distribution of well-being in Equality provides a reason to choose this outcome. On this view, total well-being is not all that matters; equality of distribution also matters. Equality, it is claimed, is an important moral consideration that the utilitarian overlooks.

Proponents of utilitarianism might respond to this objection in four ways . We will go through these in turn.

Accommodating the Intuition

The first response to this objection is to point out that utilitarians accept that equality is an extremely important guiding concept in our everyday decision making. The difference is merely that utilitarians value equality because of its instrumental benefits, rather than because it is intrinsically important.

Utilitarians care deeply about equality largely because most goods exhibit *diminishing marginal utility*. This means that the more an individual already has of a particular good, such as money or nice clothes, the less they benefit from having even more of it. This provides a strong instrumental reason for utilitarians to care about equality in distributing goods. The diminishing marginal utility of goods implies that we can often increase overall well-being by redistributing from the haves to the have-nots. Also, excessive inequality between people may cause social conflict and be bad for society in the long run. This provides an additional reason to prefer equal distributions of wealth among people.

As a practical matter, many utilitarians use their time and money to improve the lives of the most disadvantaged groups in the world, trying to create a more equitable world for all.

Debunking the Intuition

A second response to the objection is to argue that since equality is instrumentally good, perhaps we get confused into thinking it is good in and of itself. A utilitarian might argue that these moral concepts are so valuable for society that we should cultivate strong character dispositions and social norms to endorse, protect and promote them. Therefore, our intuition against particular inequitable outcomes may just result from us having embraced a general moral norm in favor of equality. While our intuitions in favor of equitable outcomes generally increase well-being, they fail when the best achievable outcome is inequitable.

Furthermore, a utilitarian might argue that our intuitions in the choice between Equality and Inequality may be unreliable since we are not used to comparing outcomes directly in the well-being of the individuals involved in them. Moral psychologist Joshua Greene argues that people find it difficult to reason quantitatively about well-being, since they confuse well-being with physical goods. Greene writes:

> We're used to quantifying stuff, things out in the world, or features of things in the world: How many apples? How much water? ... How much money? But we don't ordinarily quantify the quality of our experiences. And thus, when we imagine possible distributions of [well-being] ... it's very hard not to think of distributions of stuff, rather than distributions of experiential quality.[1]

This confusion that Greene describes may mislead our intuitions in the choice between Equality and Inequality because we are intuitively used to thinking about goods as having diminishing marginal utility. If the numbers in the hypothetical choice between Equality and Inequality represented levels of goods and not well-being, the utilitarian would choose the equal outcome. However, it would be a mistake to apply this thinking when comparing distributions of well-being, since well-being does not exhibit diminishing marginal utility. By definition, a given unit of well-being is equally valuable to the person to whom it accrues no matter how well-off they are already. So we must consider this in our intuitions.

If we specify the original example in money, and use the standard economic measure of the rate of diminishing returns to money,[2] then we should represent the choice between Equality and Inequality as

follows: In Equality 1,000 people get $13,500, while in Inequality 500 people get $10,000 and 500 people get $50,000.

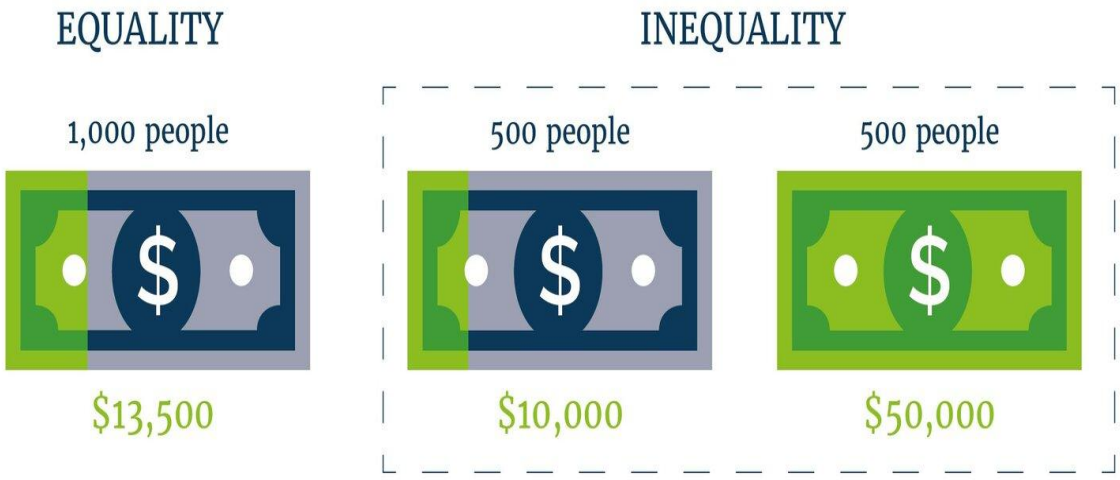

We think that, put in these terms, it is no longer intuitively obvious that the utilitarian choice, favoring Inequality over Equality, is the wrong one. Everyone would agree that some level of inequality can be justified by greater total well-being; how exactly to make this trade-off is tricky, and it is not obvious that utilitarianism gets the wrong answer.

Attacking the Alternatives

The third line of response is to *attack the available alternatives* to utilitarianism to show that they have even more counterintuitive implications. For example, egalitarianism is subject to the *Leveling Down Objection* that it countenances harming the well-off (without benefiting anyone) since that has the effect of increasing equality.[3] If you don't think that there's anything good about such leveling down then your judgment supports the utilitarian view that equality is only valuable instrumentally: valuable when and because it serves to promote overall well-being.[4]

We can also show that the alternatives to the utilitarian distribution of well-being violate a principle called *ex ante Pareto*: that, in a choice between two gambles, if everyone would rationally prefer gamble A to gamble B, then gamble A is better than gamble B.

To see this, suppose that you can choose between one of two options for distributions of well-being, *Safe Bet* and *Risky Gamble*: In Safe Bet both Abel and Beth are guaranteed 45 well-being. In Risky Gamble a fair coin is tossed. If it lands heads, Abel receives 80 well-being and Beth receives 20 well-being. If it lands tails, Beth receives 80 well-being and Abel receives 20 well-being.

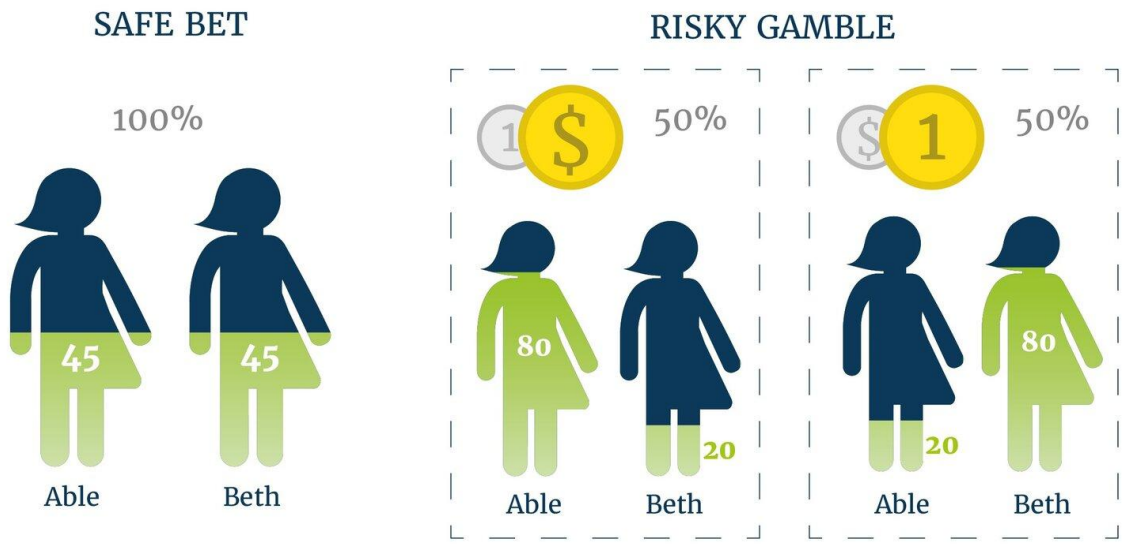

We can stipulate that both Abel and Beth would rationally prefer Risky Gamble over Safe Bet. They do so because their expected well-being in Risky Gamble is 50 (=(80+20)*50%), which is higher than the expected well-being in Safe Bet of 45.[5] So, if we follow ex ante Pareto then we should prefer Risky Gamble to Safe Bet. However, Risky Gamble results in an unequal outcome. Those who preferred *Equality* to *Inequality* must therefore also prefer *Safe Bet* to *Risky Gamble*. They must do so even though doing so is against the best interests of all parties concerned.

This objection has been developed formally by the economist John Harsanyi, who proved it in his utilitarian aggregation theorem.[6]

Tolerating the Intuition

Finally, proponents of utilitarianism may once again "bite the bullet" and simply accept that utilitarianism sometimes conflicts with our intuitions about equality. These conflicts occur because the outcomes favored by utilitarianism are those with the highest possible levels of well-being, which do not always coincide with the most egalitarian outcomes. Utilitarians will insist that it's better for people overall to be better-off (even if these benefits are not distributed equally) than for them all to share equally in a lower level of well-being.

Resources and Further Reading

- Katarzyna de Lazari-Radek & Peter Singer (2017). *Utilitarianism: A Very Short Introduction*. Oxford: Oxford University Press. Chapter 4: Objections.
- J. J. C. Smart & Bernard Williams (1973). *Utilitarianism: For and Against*. Cambridge: Cambridge University Press.