Philosophy and Phenomenological Research
Vol. LX, No. 3, May 2000

The Ethics of Belief

RICHARD FELDMAN
University of Rochester

In this paper I will address a few of the many questions that fall under the general heading of "the ethics of belief." In section I I will discuss the adequacy of what has come to be known as the "deontological conception of epistemic justification" in the light of our apparent lack of voluntary control over what we believe. In section II I'll defend an evidentialist view about what we ought to believe. And in section III I will briefly discuss apparent conflicts between epistemic considerations and moral or other considerations.

I. Epistemic Deontologism[1]

A. The Problem

Our talk about epistemic matters parallels our talk about ethical matters in noteworthy ways. In the case of ethics, we say that a person *ought* to perform a certain action, that someone *should* not do a certain thing, that people have *obligations* to act in some ways, that they have *rights* and *duties*, and that they deserve *praise* or *blame* for what they have done. We make seemingly analogous epistemic judgments about beliefs. We say that a typical well-informed contemporary American ought to believe that the Earth revolves around the Sun and should not believe that the Earth is flat. We say that a person has a right to believe one thing and perhaps a duty to refrain from believing something else. We sometimes praise those who believe the things they should and we criticize those who fail in their believings. We can describe all these judgments as *deontological judgments about beliefs*.

In "The Deontological Conception of Epistemic Justification," William Alston says that the most natural way to understand epistemic justification is deontological.[2] By this he means that it is to be understood in terms of epis-

---

1 This section is a shortened version of my paper, "Doxastic Voluntarism and the Ethics of Belief," forthcoming in a collection of papers on epistemic responsibility edited by Matthias Steup.

2 William Alston, "The Deontological Conception of Epistemic Justification," *Philosophical Perspectives* 2 (1988): 257–99. See especially p. 258. Reprinted in William Alston, *Epistemic Justification* (New York: Cornell University Press, 1989).

temic "obligation, permission, requirement, blame, and the like."[3] He regards *"requirement, prohibition,* and *permission* as the basic deontological terms" and the other terms as derivative ones.[4] Alston eventually argues that deontological conceptions are in the end ill-suited to epistemic purposes.[5] Recently, Alvin Plantinga and Alvin Goldman have independently argued that the viability of a deontological conception of epistemic justification is crucial to the debate between internalists and externalists about epistemic justification.[6] Goldman thinks that a central, but mistaken, line of support for internalist theories begins with the assumption of a deontological account of justification. Plantinga also argues that internalism derives much of its support from a deontological view of justification.[7] The merits of the deontological view of epistemic justification thus is of considerable epistemological significance to contemporary epistemology.

A central problem that both Plantinga and Alston find with deontological judgments about beliefs is that they presuppose that we have voluntary control over what we believe. Yet, reflection on our mental lives suggests that we have no such control. Alston writes of one particular deontological analysis:

> ...this conception of epistemic justification is viable only if beliefs are sufficiently under voluntary control to render such concepts as *requirement, permission, obligation, reproach,* and *blame* applicable to them. By the time honored principle that "Ought implies can", one can be obliged to do A only if one has an effective choice as to whether to do A.[8]

He goes on to argue that we don't have an effective choice over what we believe. In the process of objecting to Chisholm's views about justification, Plantinga says of a particular proposition that "...whether or not I accept it is simply not up to me; but then accepting this proposition cannot be a way in which I can fulfill my obligation to the truth, or, indeed, *any* obligation..."[9] Thus, according to Plantinga, our lack of control over beliefs implies that

---

3 "The Deontological Conception of Epistemic Justification", p. 257.

4 Ibid.

5 The most prominent advocate of this sort of analysis is Roderick Chisholm. He defends this view as recently as in his in 1991 paper, "Firth and the Ethics of Belief," *Philosophy and Phenomenological Research* 50 (1991): 119–28. For a thorough list of Chisholm's publications on this topic, see Susan Haack, "'The Ethics of Belief' Reconsidered," in Lewis E. Hahn, (ed.) *The Philosophy of Roderick Chisholm* (LaSalle: Open Court, 1997), pp. 129–44. See especially footnote 2. See also Hilary Kornblith, "Justified Belief and Epistemically Responsible Action," *Philosophical Review* 92 (1983): 33–48 for an influential discussion of the idea that a belief is epistemically justified provided it is formed in an epistemically responsible way.

6 See Alvin Plantinga, *Warrant: The Current Debate* (New York: Oxford University Press, 1993) and Alvin Goldman "Internalism Exposed," the *Journal of Philosophy* 96 (1999): 271–93.

7 *Warrant: The Current Debate*, Chapter 1.

8 "The Deontological Conception of Epistemic Justification," p. 259.

9 *Warrant: The Current Debate*, p. 38.

they are not the sort of thing that can be a matter of obligation, and this undermines Chisholm's deontological conception of epistemic justification. Matthias Steup presents a similar argument, though he goes on to defend the deontological conception on the grounds that belief is voluntary.[10]

For the purposes of the discussion that follows it will be helpful to distinguish two elements of the arguments just presented. Their target is a deontological conception of justification, a conception according to which epistemic justification is to be understood or analyzed in terms of the deontological concepts of obligation, requirement, and the like. This conception is "viable", in Alston's terms, only if belief is sufficiently under our voluntary control. Presumably, the deontological conception is viable only if it can be true that we are required to believe things, that we ought not believe other things, and so on. That is, the deontological conception of epistemic justification is viable only if deontological judgments about beliefs are sometimes true. My goal in what follows is to argue that deontological judgments are sometimes true and that good sense can be made of the idea that we can have epistemic requirements, obligations, and the like. I will not be arguing for the further claim that epistemic justification should be analyzed in deontological terms.

We can formulate the key part of the argument at issue in the following way:

The Voluntarism Argument

1. If deontological judgments about beliefs are true, then people have voluntary control over their beliefs.

2. People do not have voluntary control over their beliefs.

---

3. Deontological judgments about beliefs are not true.

Epistemologists have three kinds of response to this argument open to them: i) they can argue that we do have the requisite sort of control over our beliefs, thereby rejecting premise (2); ii) they can argue that deontological judgments do not have voluntarist implications, thereby rejecting premise (1); or iii) they can accept the argument and admit that the familiar deontological terms of epistemic appraisal really are inapplicable.[11] This in itself is a surprising conclusion, whether or not a deontological analysis of justification is acceptable.

---

[10] Matthias Steup, "The Concept of Epistemic Justification," Chapter 4 of *Contemporary Epistemology* (Upper Saddle River: Prentice-Hall, 1996).

[11] This last alternative implies that either epistemic justification is not to be analyzed in deontological terms or else that epistemic justification is itself inapplicable to our beliefs, just like the other deontological epistemic terms.

I will discuss response (i) in section B below and response (ii) in section C. I will not discuss response (iii) except in passing and by implication.

B. Doxastic Voluntarism

In this section I will examine some arguments for doxastic voluntarism and some arguments against it. I will eventually defend a version of doxastic voluntarism and I will also argue that this fact is of absolutely no epistemological significance and that it does nothing to help resolve the voluntarism puzzle. Although some philosophers have argued that it is a conceptual impossibility that anyone form a belief voluntarily, I will not discuss that view.[12] Instead, I will examine the claim that it is a contingent matter of fact that we lack that ability.

B1. The Argument for Involuntarism

Alston has given the most thorough defense of the contingent inability thesis, the thesis that as a contingent matter of fact people are not able to acquire beliefs voluntarily. Alston's paper includes an excellent survey of a variety of notions of voluntary control. For each type, he argues that we lack that sort of control over beliefs. (Alston does admit that there is one very weak notion of control that does apply to belief. But he contends that this sort of control does not provide an adequate basis for a good response to the Voluntarism Argument.)

Alston begins by discussing *basic voluntary control.*[13] We have basic voluntary control over those actions that we can "just do." Simple bodily motions are the prime examples. I can just raise my hand, close my eyes, and bend my knee. Some people, but not I, can wriggle their ears and curl their tongue. Alston correctly says that forming a belief is not like that. We can't just do it at will.

Alston turns next to *nonbasic immediate voluntary control.*[14] One has this sort of control over the things one can do right away by doing something else, typically something over which one has basic voluntary control. Standard examples are opening doors and turning on lights. We can, in typical circumstances, do these things simply by moving our bodies in the appropriate ways. There's vagueness here concerning what counts as "right away" but

---

12 For defenses of the conceptual impossibility claim, see Bernard Williams, "Deciding to Believe," *Problems of the Self* (New York: Cambridge University Press, 1973), Jonathan Bennett, "Why Is Belief Involuntary?," *Analysis* 50 (1990) and Dion Scott-Kakures, "On Belief and the Captivity of the Will," *Philosophy and Phenomenological Research* 53 (1993). For discussion of Williams, see Barbara Winters, "Believing at Will," *Journal of Philosophy* 76 (1979). For a discussion of Scott-Kakures' article, see Dana Radcliffe, "Scott-Kakures on Believing at Will," *Philosophy and Phenomenological Research* 57, 1 (1997).

13 "The Deontological Conception of Epistemic Justification," Section III, pp. 263–68.

14 See Section IV, pp. 268–74.

that vagueness is in no way problematic. This is because the boundary between nonbasic immediate voluntary control and the next weaker kind of control, *long range voluntary control*, is acceptably imprecise.[15] The sorts of things over which we have long range voluntary control are the sorts of things we can do over time by doing other things. Perhaps painting my house is an example. Or, more precisely, I have long range voluntary control over what color my house is because I can do things like paint it. Finally, there is *indirect voluntary influence*.[16] This is the kind of control we have when we can undertake a course of action that may affect some condition over the long term. Perhaps a person has indirect voluntary influence over the condition of her heart, since diet and exercise (courses of action we can more directly control) can affect it.

Alston claims that we have only this weakest sort of control over our beliefs. One can engage in a course of action for the purpose of inculcating or losing a particular belief. Consider my belief that the Earth is not flat. There's nothing I can do to get myself to stop believing that right away. It's not like shutting the door or turning out the lights. And there's nothing much I can do long range to control it either. I might enroll in the Flat Earth Society, read conspiracy literature asserting that satellite photos are all phony, and so on. Perhaps this will help rid me of my belief. Alston would agree that it might. But this gets us, at most, indirect voluntary influence, and this is not the sort of effective voluntary control required to refute The Voluntarism Argument.[17]

I believe that, for most part, what Alston says is right. However, I will argue that we have considerably more control over some of our beliefs than Alston acknowledges. Still, this control does not undermine the basic idea behind The Voluntarism Argument, although it may show that the argument needs reformulation.

B2. An Argument for Doxastic Voluntarism

My argument for doxastic voluntarism begins with the assumption that there are states of the world over which people have nonbasic voluntary control. For example, I have nonbasic voluntary control over whether the lights in my office are on. All I have to do is move in a certain way to get the lights on or off. And I can do this. The next step of the argument notes that my belief about whether the lights are on tracks their actual state almost perfectly. As a result, I have a similar amount of control over whether I believe that the lights are on. All I have to do is move a certain way and then

---

15 See Section V, pp. 274–77.

16 See Section VI, pp. 277–83.

17 Perhaps this point suggests that both (2) and the consequent of (1) should be about *effective* voluntary control.

I'll have the relevant belief. More generally, when I have control over a state of the world and my beliefs about that state track that state, then I have just as much control over my belief about the state as I have over the state itself. Thus, we have nonbasic immediate voluntary control over our beliefs about states of the world over which we have control, provided our beliefs are responsive to those states. Furthermore, if we know that we will respond in some mistaken way to some state of the world over which we have control, we also have control over the resulting (erroneous) belief. So, doxastic voluntarism is true after all. Premise (2) of The Voluntarism Argument is false.

I believe that the existence of nonbasic voluntary control over beliefs can have prudential and moral significance. If the department chair announces that she'll give a raise to all and only those members of the department who in 30 seconds believe that the lights in their office are on, I'll head for the light switch and turn on the lights to make sure I have the desired belief. If the chair perversely announces that the graduate students will be mercilessly tortured—say, by being forced to take additional prelims—unless in 30 seconds I believe that my lights are on, then I'd better make sure that I have that belief. I'm in control.

Thus, we do have control over many of our beliefs. Of course, we don't often exercise this control. That is, we don't often believe things voluntarily. And this leaves this defense of doxastic voluntarism without a great deal epistemological significance. The central reason for this is that we make deontological epistemic judgments about beliefs that we can't control and these judgments are as routine and commonplace as are judgments about beliefs that we can control. Thus, our ability to control what we believe in the way described here is epistemically insignificant. We can take this into account by reformulating The Voluntarism Argument:

The Voluntarism Argument (Revised)

1. If deontological judgments about beliefs concerning states of the world people can't control are true, then people have (effective) voluntary control over those beliefs.

2. People do not have (effective) voluntary control over those beliefs.

---

3. Deontological judgments about those beliefs concerning are not true.

This argument is as troubling as the original. It implies that an enormous number of the deontological epistemic judgments we routinely make cannot be true. Furthermore, the argument might be extended to beliefs that we can control but have not formed in this voluntary way. Thus, I don't think that

the fact that we have nonbasic immediate voluntary control over some beliefs provides the basis for an effective defense of epistemic deontologism.[18]

B3. A More Robust From of Doxastic Voluntarism

Compatibilists in the free will debate contend that we voluntarily perform an action when the action has the right sort of cause. A defender of doxastic voluntarism can argue that, analogously, we believe voluntarily when we believe as a result of the right sort of causal process. Roughly, the idea is that when unconstrained deliberation about evidence leads to belief, we believe voluntarily.[19] This is an interesting line of thought, and it is difficult to refute because we lack a fully adequate understanding of what counts as "the right sort of causal process." Still, I think that there are good grounds to reject the idea that the deliberative process that typically leads to belief is the kind of process that makes its product a voluntary action.

A key fact about the clearest cases of voluntary action is that the action is caused by an intention to perform that action. I turn on the lights because I intend to turn on the lights; I type a certain word because I intend to type that word. No doubt there are puzzles about the details of these matters, such as puzzles resulting from the fact that our intentions are not as precise as the resulting actions—I need not have intended to turn on the light in exactly the way I did. But however we deal with these details, the case of belief is dramatically different. We simply don't, in the typical case, form a belief as a result of an intention to form that belief. It may be that in some sense we do control what we believe, in much the way we control other involuntary processes in us. However, epistemic deliberation does not result in effective intentions to believe. Except in rare cases, we don't form intentions to believe. But such intentions are essential to voluntary control.[20] The cases of voluntary belief formation that I described earlier make it easy to see what it really would be like to believe voluntarily. In those cases, I do believe as a result of an intention to believe. But those are, for this very reason, quite unlike ordinary cases.

Furthermore, if this defense of doxastic voluntarism were sound, then it seems that many other plainly involuntary behaviors would also be voluntary. Deliberating about something can result in states other than belief. It

---

18 If this revised argument is strong, then it would also be reasonable to reject deontological analyses of epistemic justification as well.

19 For a defense of this view, see Matthias Steup, "Doxastic Voluntarism and Epistemic Deontology," forthcoming. For a similar view, see James Montmarquet, "The Voluntariness of Belief," *Analysis* 46 (1986). Bruce Russell endorses Steup's view in "Epistemic and Moral Duty," forthcoming.

20 Perhaps the process in which by which we formulate intentions and act on their basis can become automated in certain ways, so that voluntary actions need not involve explicit conscious formulations of intentions. Belief formation is not like that either. It does not typically involve the formation of intentions to believe at all.

can result in desires or in panic. If we say that belief is voluntary because it is the outcome of deliberation about evidence, then it is hard to see why we shouldn't say that these other outcomes of deliberation are also voluntary. Yet it is clear that they are not. Again, what seems clearly missing is the right sort of intention.

Thus, I reject this more robust form of doxastic voluntarism. There is good reason to say that the deterministic processes that typically lead to action render those actions voluntary. The processes lead to effective intentions to act. Deliberating about evidence does lead to belief, but not via such intentions. Such deliberation does not make belief voluntary. We need a better response to The Voluntarism Argument.

C. Ought and Can

The second main sort of response to The Voluntarism Argument denies that the deontological judgments about beliefs imply that those beliefs are voluntarily adopted. It denies premise (1) of the argument. It may be that there are differences among the various assertions that I've described as "deontological." I'll focus first on judgments about what one is obligated to believe and what one ought to believe and then discuss other judgments later in the section.

One way to defend deontological epistemic judgments in the light of doxastic involuntarism is to argue that we can have epistemic obligations even though we can't fulfill them (or even if we can't help but fulfill them). This is to deny that "epistemically obligated to" or "epistemically ought to" implies "can". And one way to make this denial plausible to show that there are other kinds of ought statements that don't imply voluntary control. I will consider several candidates for this other kind of ought.

C1. Contractual Obligations

You can have an obligation to pay your mortgage even if you don't have the money to do so. Perhaps students in a class have an obligation to do the course work even if they are incapable of doing it. Other examples of this sort are rather easy to construct. Perhaps epistemic obligations are analogous to these financial and academic obligations. Though I once defended this view,[21] it now strikes me as an implausible model for epistemic obligations. The obligation to pay one's mortgage and the obligation to do one's course work are contractual obligations, although in the latter case the contract is in some sense implicit. It's difficult to see any basis for saying that we all have some sort of contractual obligation to believe things. Surely no such contract is explicit, and nothing analogous to enrolling in a course establishes an implicit contract.

---

21 "Epistemic Obligations," *Philosophical Perspectives* 2 (1988), pp. 240–43.

C2. Paradigm Obligations

In a recent paper, Nicholas Wolterstorff says that there are two kinds of obligations, *paradigm* obligations and *responsibility* obligations.[22] Only obligations of the latter sort are associated with voluntariness. As examples of paradigm obligations, he presents:

1. "You ought to be walking on it in two weeks"—said by a physician as he finishes binding up a person's sprained ankle.

2. "That's strange; you ought to be seeing double"- said by a psychologist to his subject while conducting an experiment in perception.

Wolterstorff suggests that epistemic obligations are similar to the obligations described by these sentences. They lack any implication of voluntary control.

No friend of epistemic deontologism should be comforted by the idea that epistemic obligations are like the obligations described by (1)–(2). This is because there are no obligations described by (1)–(2). Sentences (1)–(2) are "ought" sentences; they are not obligation sentences and they cannot be paraphrased in any straightforward way into obligation sentences. Your ankle has no obligation of any sort to heal; you have no "perceptual obligation" to see double. So, if there are epistemic obligations, they are not like the obligations described here, since there are no obligations described here.

Furthermore, the ought sentences Wolterstorff describes are not relevantly like epistemic oughts. (1) and (2) describe normal, or paradigmatic, behavior. Thus, for example, baring unforeseen developments, your ankle will heal in 2 weeks. But I don't think that this carries over to the epistemic case. That is, epistemic oughts don't describe paradigmatic or normal function. Some researchers report that people typically make various unjustified inferences and predictably form unreasonable or erroneous beliefs.[23] Whatever the proper interpretation of this research actually is, it is at least possible that people normally make epistemic errors. It may be that we epistemically ought not do what we normally do. Epistemic oughts are not paradigm oughts.[24]

---

22 Nicholas Wolterstorff, "Obligations of Belief: Two Concepts," in *The Philosophy of Roderick Chisholm*, Hahn, Lewis E (LaSalle, Illinois: Open Court, 1997).

23 For a discussion of many of the examples allegedly showing that people are irrational and an examination of their philosophical implications, see Edward Stein, *Without Good Reason: The Rationality Debate in Philosophy and Cognitive Science* (Oxford: Clarendon Press, 1996).

24 Thus, a researcher might say, while awaiting the subject's reply to a question that nearly everyone misses, "He ought to make the wrong inference here." The use of "ought" here is Wolterstorff's paradigm ought. But to describe the inference as "wrong" is to say that it is not the case that the person epistemically ought to do what the researcher expects him to do.

C3. Role Oughts

There are oughts that result from one's playing a certain role or having a certain position. Teachers ought to explain things clearly. Parents ought to take care of their kids. Cyclists ought to move in various ways. Incompetent teachers, incapable parents, and untrained cyclists may be unable to do what they ought to do. Similarly, I'd say, forming beliefs is something people do. That is, we form beliefs in response to our experiences in the world. Anyone engaged in this activity ought to do it right. In my view, what they ought to do is to follow their evidence (rather than wishes or fears). I suggest that epistemic oughts are of this sort—they describe the right way to play a certain role. Unlike Wolterstorff's paradigm oughts, these oughts are not based on what's normal or expected. They are based on what's good performance. Furthermore, it is plausible to say that the role of a believer is not one that we have any real choice about taking on. It is our plight to be believers. We ought to do it right. It doesn't matter that in some cases we are unable to do so. Thus, I reject the first premise of the revised Voluntarism Argument. Even in cases in which a believer has no control at all, it makes sense to speak of what he ought to believe and ought not believe.

I think that most, possibly all, of the deontological terms we ordinarily use to evaluate beliefs can be explained in similar terms, although the case for some may be weaker than the case for others. Thus, if the standards of good believing allow either of two attitudes in a given situation, then it will be true that the believer has a *right* to either of those attitudes and that either of them is *permitted*.[25] Other attitudes may be *prohibited*. It seems to me reasonable to say that when only one attitude is permitted, then one has an *epistemic obligation* to have that attitude. (I will add a qualification to this claim later in the paper.) It may be that some terms, especially those associated with *praise* and *blame*, are to be reserved for voluntary behavior. Even here, the case is less than perfectly clear, since we do praise and blame people for attributes, such as beauty, that they are unable to control. Thus, I conclude that deontological epistemic judgments can be true even if doxastic voluntarism is false.

C4. Is This Deontology?

I want to contrast the view I'm defending from a different view about epistemic deontologism. Possibly some critics object primarily to the idea that to be epistemically justified in a belief is to do the best that one can do with respect to that belief. Their target really does imply voluntarism. I haven't attempted to defend that view, and I don't think that it is true. I've taken on the much more modest task of defending the legitimacy of the widespread use of deontological language about belief. My contention is that we can have

[25] I'll return to some related issues in section II, part A.

epistemic requirements, permissions, and the like even if voluntarism is false.

In his influential paper on this topic, Alston begins by saying that it is "natural" to understand epistemological terms in a "'deontological' way, as having to do with obligation, permission, requirement, blame, and the like."[26] As an example of the view he wants to question, he mentions Carl Ginet's view that explains justification in terms of whether one is as confident as one ought to be.[27] But by the end of his paper Alston seems to be directing his attack more narrowly. In his concluding paragraph he characterizes epistemic deontologism as the view that analyzes "epistemic justification in terms of freedom from blame for believing."[28] I haven't attempted to defend this sort of deontologism here. Moreover, I don't think that this more narrowly defined sort of deontologism is so natural or common. My defense is only of the possible truth of the deontological epistemic claims with which Alston began his paper, claims to the effect that people can have epistemic rights, duties, permissions, etc. They can, no matter what the truth about doxastic voluntarism is.[29]

Given that deontological epistemic judgments can be true, I turn next to a view about the conditions under which they are true.

II. Epistemic Oughts

A. Evidentialism

In this section I will defend a variation on William K. Clifford's frequently quoted claim that "It is wrong, always, everywhere, and for anyone to believe anything upon insufficient evidence."[30] It is not my purpose here to explain or defend the thesis Clifford asserted with these words. My view may differ from Clifford's in any or all of the following ways. First, he may have been making a *moral* evaluation of believing on insufficient evidence.[31] The claim

---

26 "The Deontological Conception of Epistemic Justification," p. 257.

27 Ibid., p. 259.

28 Ibid., p. 294.

29 It is possible that some critics will contend that sentences like "Smith ought to believe that the earth is not flat" are ambiguous between deontological and non-deontological senses. They might then object that I've merely argued for the possible truth of the non-deontological interpretation of the sentence. I don't see any reason to admit that there are multiple senses of this sort. The deontological sense of the sentence just mentioned seems to be nothing other than the conjunction of the non-deontological sense with the proposition that it is up to Smith whether Smith does believe the earth is not flat. Obviously, so interpreted, it does imply voluntarism. But the question with which we began was whether the original claim was compatible with voluntarism. I don't think that this compatibility can be ruled out by fiat.

30 W. K. Clifford, "The Ethics of Belief," reprinted in *The Theory of Knowledge* 2nd Edition, edited by Louis P. Pojman, (Belmont, California: Wadsworth, 1999), pp. 551–54.

31 In "The Ethics of Belief Reconsidered" Susan Haack argues that Clifford did not distinguish epistemic and moral senses of the key terms.

I want to make concerns epistemic rather than moral evaluation. Second, Clifford evaluates believing on insufficient evidence as *wrong* whereas the claim I will defend is that one ought not believe on insufficient evidence. Whether there is a difference between a thing being wrong and being something one ought not do is unclear to me. If there is a difference, it is the latter claim that I want to defend. Third, in the quoted passage Clifford only objects to believing on insufficient evidence. He does not say that one ought to believe when one does have sufficient evidence. However, I do want to make this additional claim. A succinct way of stating my thesis is that one always ought to follow one's evidence. This *evidentialist* thesis is what I will defend.[32]

The following principle provides a preliminary statement of the evidentialist idea:

O1. For any proposition p, time t, and person S, S epistemically ought to have at t the attitude toward p that is supported by S's evidence at t.

The attitude in question can be belief, disbelief, or suspension of judgment.[33] It's important to realize that (O1) does not imply that a person ought to believe a proposition simply because some of the person's evidence supports that proposition. Belief is the attitude to have only if the evidence on balance supports p.

Some philosophers who are generally sympathetic to evidentialism may think that merely having one's evidence on balance support a proposition is insufficient for it to be true that one epistemically ought to believe the proposition. That evidential support might be exceedingly weak, and some will think that stronger evidence is required. Those who insist on such a condition can interpret (O1) in a way that fits with this idea. On this interpretation, believing p is supported by one's evidence only if the evidence supports p sufficiently well.

There is, however, reason to doubt (O1). Suppose that a person has evidence that conclusively establishes some proposition, q. There are then a huge number, perhaps an infinite number, of obvious logical consequences of q that are also supported by this evidence. For example, the evidence supports q disjoined with every other proposition, or at least every other proposition S understands. (O1) seems to imply that S ought to believe all these disjunctions. Many of these disjunctions will be trivial; many will be of no practical

32 In "Evidentialism," *Philosophical Studies* 48 (1985): 15–34, Earl Conee and I defend the evidentialist thesis that believing a proposition is justified if and only if believing the proposition "fits" one's evidence.

33 Those who think that beliefs come in degrees may wish to replace this principle by one saying that one ought to believe propositions to the degree that they are supported by their evidence.

or theoretical interest. It may be that in some sense believing these propositions is justified—the person has good support for them—but it is not true that he *ought* to believe them. To do so would constitute an exceedingly foolish use of S's cognitive resources.[34]

Whether this objection is a good one depends in part upon what exactly goes into believing a proposition. The more dispositional belief is, the less that (O1) demands of us. If the person with conclusive evidence for q can satisfy its demands by being disposed to assent to each of the disjunctions if asked about them, then (O1) doesn't really require that he expend any cognitive resources. He doesn't have to do anything to meet its standards. If, on the other hand, it takes more than that to believe a proposition, then perhaps (O1) is excessive. I'm inclined to think that (O1) does ask too much of us, but I won't argue for that here. Instead, I propose the following refinement that avoids the problem:

O2. For any person S, time t, and proposition p, if S has any doxastic attitude at all toward p at t and S's evidence at t supports p, then S epistemically ought to have the attitude toward p supported by S's evidence at t.

(O2) in effect conjoins three principles: if a person is going to adopt any attitude toward a proposition, then that person ought to believe it if his current evidence supports it, disbelieve it if his current evidence is against it, and suspend judgment about it if his evidence is neutral (or close to neutral). A person might never even consider a proposition and not be about to adopt any attitude toward it. In that case, (O2) does not imply that there is any attitude the person ought to have toward the proposition. Thus, (O2) avoids the problem concerning the use of cognitive resources that (O1) faced. (O2) merely states which attitude to take toward propositions about which you are going to have one attitude or another. It does not instruct you to believe logical consequences that you don't entertain or to otherwise squander precious resources.

B. Requirements and Permissions

(O2) helps clear up a confusing issue about whether epistemic principles state what you are required (ought) to believe or what you are permitted to believe. Suppose some unimportant and unconsidered proposition obviously follows from your evidence. For the reasons just stated, this fact conjoined with (O2) doesn't imply that you ought to believe it. But (O2) does imply that if you take any attitude toward the proposition, it ought to be belief. So, it's not

[34] Alvin Goldman makes this point in "Epistemics: The Regulative Theory of Cognition," *The Journal of Philosophy* 75 (1978): 509–23.

just that believing it is permitted. That's the only permissible attitude. It's the one you ought to have if you have any attitude. Thus, no attitude is epistemically required, but only one is epistemically permitted.

(O2) thus rules out the possibility that each of two attitudes toward a proposition is permitted at a time. Some critics might regard this as a mistake. To see whether it is, let's begin by considering a moral analogue of this possibility. On standard views, there are cases of ties in morality: each of two actions is morally permissible and neither is obligatory. It's permissible for me to give my extra money to the Cancer Society and it's permissible to give it to the Heart Association. It's not permissible to spend it frivolously. It follows then, that it is permissible to give my money to the Heart Association and also permissible to refrain from giving my money to the Heart Association. Each of two incompatible actions is thus permissible. This is a consequence of the existence of moral ties. If the epistemic case is similar, then there are situations in which each of two attitudes toward one proposition is permissible. But I think that in the epistemic world, things aren't quite the same. There are no ties, in the relevant sense.

Potential cases of epistemic ties are of three sorts. The first, and easiest to dispose of, are cases in which the evidence for and against a proposition is equally weighty. In that case, evidentialism implies that the sole acceptable attitude is suspending judgment. Neither belief nor disbelief is permitted. That strikes me as the right result.

The second kind of case consists in those in which two competing propositions are equally well supported by the evidence. For example, there is equally good evidence supporting the guilt of each of two suspects for a crime that could have been committed by only one person. Casual reflection might suggest that you could believe that either is guilty, but obviously, not both. In that case, it is permissible to believe that suspect A is guilty and suspect B is not, and it is permissible to believe that suspect B is guilty and A is not. In fact, however, I think that in the situation described you should not believe that either suspect is guilty. Again, if you have an attitude toward these propositions, you should suspend judgment about each, though perhaps not about their disjunction.

There are possible cases like the one just described in which some action is forced on you and it seems that you ought to form a belief about which action to take. Suppose that you know that one of two boxes in front of you has a bomb that is about to explode. You have time to open only one box and disarm the bomb if it's there. You have no reason to prefer one box to the other. Surely you ought not suspend judgment about the location of the bomb and do nothing. That evaluation is correct, but it's correct because you ought not do nothing. Failing to act would be imprudent, perhaps immoral. But you needn't have a belief about the location of the bomb in order to act. You can simply choose to open one box. And if you are the sort of person

who can only pick a box in this case if you have the relevant belief, then there may well be a sense in which you ought to form a belief about where the bomb is (if you can). But the ought here is a prudential or moral ought, not an epistemic ought.[35] You epistemically ought to suspend judgment.

The third and most difficult sort of case to think about is one in which you have some modest amount of evidence supporting a proposition. It's tempting to think that it's permissible to believe the proposition and it's also permissible to suspend judgment about it, on the grounds that belief requires more than just this modest amount of evidence. In support of this line of thought is the idea that believing on only modest amounts of evidence involves taking some epistemic risk. There is no unique amount of risk that is right or rational. Rather, people simply have varying attitudes toward this risk. My own inclination, which I won't defend here, is that you should believe when your evidence is supportive rather than neutral, even if the evidence is not at all decisive. Those who say that one's evidence supports believing a proposition only when that evidence is sufficiently strong will in all probability be faced with borderline cases. In a case in which this modest amount of evidence is slightly in favor of a proposition, they might think that believing it is permissible and that suspending judgment is permissible. Of course, they are not forced to this conclusion, since they might hold that it is simply indeterminate which attitude one is permitted to have and that it is not simply true that either attitude is permitted.[36]

The evidentialist account of what we ought to believe relies crucially on the notion of evidential support. Analyzing this notion in a fully satisfactory way is no easy task. Among the problems to be worked out is that of determining which logical consequences of a body of evidence are supported by that evidence. There are possible cases in which a person has evidence that implies some proposition, but the connection between that evidence and that consequence is distant and difficult to see. It may be well beyond the intellectual talents of the person. I believe that in such cases the person ought not believe the consequence. Given his failure to see that it is a consequence, to believe it (barring other reasons to believe it) would be rash. Furthermore, as I understand (O2), it has exactly the right result in this sort of case. The fact that a person's evidence implies some proposition is not sufficient for the evidence to provide evidential support for the proposition. Roughly, only those propositions whose connection to the evidence the person apprehends

---

35 Perhaps the reason it seems right to say that "You ought to form the belief *if you can*" is that the "ought" in this judgment is one that does imply "can". This is further evidence that the intended judgment here is not an epistemic judgment.

36 Hud Hudson helped me to see this point. Notice that philosophers who think that ties are possible can still be evidentialists. That is, they can say that you always follow your evidence and never believe contrary to your evidence. They just think that there are cases in which the evidence leaves open two possibilities and there's nothing else that could eliminate one.

are actually supported by his evidence. And I think ascertaining this connection is itself an element of the person's evidence. These issues are complex, however, and I will not pursue them further here.

This completes my description of the evidentialist idea that one epistemically ought always to follow one's evidence. I turn now to a defense of this idea.

C. Evidentialism and Epistemic Value

One way to explain why we ought to do something is to show that it is a means to some goal that we have. A similar, but slightly different, way is to show that it is a means to some valuable result. Both alternatives explain what we ought to do instrumentally. To use the former alternative to show that we ought to follow our evidence would require showing that people have some goal—true belief or knowledge, perhaps—and that following their evidence is the proper means toward that goal. If the oughts in question are supposed to be means to goals that people actually have, then it seems that only people who do have the epistemic goals just mentioned would be subject to the relevant epistemic requirements. However, (O2) is not restricted in that way. It says that all people epistemically ought to follow their evidence, not just those who have adopted some specifically epistemic goals.[37]

One might argue that all people naturally have epistemic goals or perhaps that the goal of true belief or knowledge is in some sense an "ideal" goal. I won't pursue those claims here. Instead, I want to defend (O2) by arguing that following one's evidence is the proper way to achieve something of epistemic value. My approach does not depend upon the assumption that epistemic value is a kind of instrumental value. Of course, the view I'll defend does not imply that beliefs having epistemic value lack any sort of instrumental value.

C1. True Belief and Epistemic Value

The dominant source in the literature for discussions of epistemic goals and epistemic value is William James' famous passage:

> There are two ways of looking at our duty in the matter of opinion,—ways entirely different, yet ways about whose difference the theory of knowledge seems hitherto to have shown little concern. We *must know the truth*; and *we must avoid error*,—these are our first and great commandments as would-be knowers...[38]

---

37 One could take epistemic oughts to be conditional. To say that one epistemically ought to believe something is to say that if one has epistemic goals, then one ought to believe it.

38 William James, "The Will to Believe," reprinted in *The Theory of Knowledge*, 2nd Edition, pp. 555–62.

One idea commonly extracted from this is that our epistemic goal is to believe all and only truths.[39] But if this implies that all people actually have the goal of believing all truths, then I doubt that it is true. People's goals are a varied lot, and I doubt that all people have this one. Moreover, believing all truths is obviously an unattainable goal. We simply can't believe all the truths. Furthermore, attaining it is not desirable. As anyone who lived through President Clinton's impeachment trial knows, there are many truths one would prefer not believing (or even considering). So, it is doubtful that believing all truths is an actual, attainable, or desirable goal.

Whatever our goals, it makes sense to suppose that believing truths (or some related state) has epistemic value, that it is a good thing from an epistemic point of view. I turn now to the idea that having true beliefs and avoiding false beliefs has epistemic value and we should follow our evidence because it is the best way to have valuable beliefs. Consider, then, the following principle:

V1. Each person maximizes epistemic value by making it the case that for every proposition p, he or she believes p if p is true and does not believe p if p is false.

One might then argue that people can best maximize epistemic value by following their evidence.

There are a few problems with a defense of (O2) that is based on (V1). For one thing, (V1) would seem better suited to a defense of (O1), since (V1) assigns value to believing all truths, not just to the ones about which one has some attitude or other. Furthermore, the alleged connection between (V1) and (O2) is based on the assumption that following your evidence is the best way to get at the truth. And I just don't see why this is true. If you are in unfortunate circumstances in which the information you have will lead you to falsehoods, following your evidence is *not* the best way to the truth. That is, it is not the most effective way to get at the truth. In these circumstances, ignoring your evidence is a better way to believe (some) truths.

A possible reply to at least the second problem just mentioned for the argument from (V1) to (O2) relies on the claim that while following one's evidence might not be an infallible way to achieve the goal stated in (V1) or even always the most effective way, it is nevertheless the best or most reasonable way to try to have true beliefs. This defense of (O2) is quite similar to the one I will present in section (C3) below and I will discuss it there.

---

[39] The idea that this is our epistemic goal is widespread. It is the starting point for numerous discussions of a variety of epistemological issues. See, for example, Alston's "The Deontological Conception of Justification," p. 258 and Richard Foley's *The Theory of Epistemic Rationality* (Cambridge, Massachusetts: Harvard University Press, 1987), p. 8.

C2. Knowledge and Epistemic Value

What James actually said in his famous passage wasn't that you should believe truths or avoid believing falsehoods. Nor did he say that you should have true belief or the avoidance of false belief as a goal. What he actually said is that our commandment, as would be *knowers*, is to *know* the truth and to avoid error. One might extract from this the idea that what has epistemic value is not mere true belief but rather knowledge. Here's one way to spell out this idea:

V2. Each person maximizes epistemic value by making it the case that for every proposition p, he or she knows p if p is true and knows ~p if p is false.

(V2), like (V1), seems better suited to support something like (O1) than (O2), but I will ignore this point for now. One argument from (V2) to (O2) goes as follows. Given that one can have knowledge only if one has good evidence for what one believes, one can get beliefs with epistemic value only by following one's evidence. In other words, following one's evidence is a necessary condition for getting what (V2) says has epistemic value. If we epistemically ought to do whatever is necessary to obtain this epistemic value, then we epistemically ought to follow our evidence.

The idea that knowledge is what has epistemic value has a certain plausibility. Knowledge is valuable, and it makes sense to think that knowledge is a particularly *epistemic* kind of value. However, the argument from (V2) to (O2) is unsound. The central problem with the argument turns on cases in which a person has strong evidence for a false proposition, *f*. According to the evidentialist position expressed in (O2), such a person should believe that falsehood. And the argument for (O2) based on (V2) relies on the assumption that the person should adopt the attitude that is necessary for him to have in order to have knowledge. But since *f* is false, there's an equally good argument for the conclusion that the person ought to believe ~*f*, since knowledge also requires truth. If adopting an attitude that will yield knowledge is what's valuable, then in such a situation adopting one attitude—believing *f*—will satisfy one necessary condition for knowledge and adopting a different attitude—believing ~*f*—will satisfy a different necessary condition for knowledge. (V2) provides no reason to think that one ought to believe *f* in this case. Thus, it fails to support (O2).

A related point concerns cases in which one has strong evidence for a proposition, but evidence insufficiently strong to yield knowledge. (O2) implies that one ought to believe in such a case. But it can't be that one ought to believe in such a case because that is a means to knowledge. The situation precludes having knowledge, yet (O2) has implications regarding what one ought to do. Again, (V2) fails to provide support for (O2).

C3. Reasonable Belief and Epistemic Value

While true beliefs may have considerable instrumental value, a person who irrationally believes a lot of truths is not doing well epistemically. In contrast, a person who forms a lot of rational but false beliefs is doing well epistemically. While knowledge also has a kind of value, seeing it as the only thing of epistemic value fails to explain what is valuable about forming beliefs that fall short of knowledge. We avoid the problems associated with identifying epistemic value with true belief or with knowledge if instead we say that what has epistemic value are rational beliefs. To do well as a believer, to achieve a kind of epistemic excellence, one must form only rational beliefs. I will discuss a defense of (O2) based on this account of epistemic value in this section.

One way to understand this evidentialist perspective on epistemic value and epistemic oughts is as follows. Consider a person who is contemplating a particular proposition. To carry out the role of being a believer in an epistemically good way, in a way that maximizes epistemic value, the person must adopt a rational attitude toward this proposition. There are other values that beliefs might yield. Some beliefs might have prudential or moral value. They might make people feel good or provide comfort for others. They might help one to undertake risky but beneficial behavior. They might give one self-confidence that can be help to advance one's career. But beliefs that are beneficial in these ways can nevertheless fail to be rational. They can lack epistemic value. To achieve epistemic value one must, in each case, follow one's evidence.

Here, then, is a principle about epistemic value that supports (O2):

V3. When adopting (or maintaining) an attitude toward a proposition, p, a person maximizes epistemic value by adopting (or maintaining) a rational attitude toward p.

Given that (V3) specifies what maximizes epistemic value, a believer epistemically ought to form attitudes as directed by (O2): he ought to follow his evidence. This defense of (O2) depends on the substantive epistemological thesis that rationality consists in making one's beliefs conform to one's evidence. This thesis is apparently denied by some, for example some reliabilists. If reliabilism has substantive implications that differ from those of evidentialism, it must be because in some cases it implies that it is rational to form beliefs in the absence of evidential support. I will not undertake a defense of this evidentialist claim here, nor will I discuss the reliabilist alternative.

Anyone who is about to adopt an attitude toward a proposition, p, and who adheres to the dictates of (O2) will get knowledge (of p or of ~p) whenever such knowledge is attainable. If the available evidence is strong enough

to yield a belief well-enough justified for knowledge, and the other conditions for knowledge are satisfied, then the person will have knowledge. Where knowledge is not attainable, the person will have reasonable belief. He'll be doing as well as he can, epistemically speaking. By seeking rational belief, then, one will get knowledge where one can. And, of course, unless one is in unfortunate circumstances in which one's evidence frequently leads to false beliefs, anyone who follows (O2) will mostly have true beliefs.

(V3) and (O2) concern themselves with how to maximize epistemic value when adopting an attitude toward a proposition. There are many ways in which one might behave over the long term that will help one to gain knowledge of important facts. The relevant behavior might include enrolling in suitable courses of study, maintaining a healthy life-style to keep one's mind sharp, and cultivating sound inferential habits. Evidentialism is silent about such practices. It focuses on the epistemic value to be obtained immediately from the adoption of an attitude toward a proposition. The way to do that, in every case, is to follow the evidence one has.

At the end of section (C1) I mentioned the idea that one might defend (O2) on the grounds that having true beliefs is epistemically valuable and that following one's evidence is the most reasonable way to try to get beliefs that have epistemic value. The current view is that reasonable beliefs are epistemically valuable and that following one's evidence is a perfect means to getting valuable beliefs. There are then two ways to support (O2). The differences between these defenses of (O2) may not be great, but there are two points worthy of note. They can be brought out by considering a person whose evidence supports a great many falsehoods. First, suppose the person follows his evidence. According to both views, this person is believing as he ought. But according to the earlier view his beliefs, being false, lack epistemic value whereas according to the current view they are epistemically valuable. Second, suppose that the person does not follow his evidence. In that case, both views imply that he is not believing as he ought, but the earlier view implies that he is, by luck, achieving epistemically valuable beliefs. In both cases, I find the implications of current view more appealing. I don't see anything epistemically good about the person who irrationally gets true beliefs. I don't think that it would be correct to say of him that he's achieved epistemic excellence, even though he's done it in an irrational way or merely by luck. Rather, I think he's failed epistemically, not only because he isn't believing as he ought but because he does not have rational beliefs. Of course, there may be some instrumental value in those true beliefs. They may help the person negotiate the world in a better way. But that is a different matter.

D. Three Objections to Evidentialism

D1. Evidence One Should Have Had

Consider a person who is negligent about collecting evidence. Suppose that I have the firm belief that

> G. Taking ginkgo supplements is a safe and effective way to improve my memory.

I have a modest amount of evidence supporting (G). I then see on the cover of a reputable magazine that it contains a major article about the merits of gingko. Rather than read the article, I avoid it for fear that it will undermine a belief I prefer to keep. Thus, I am negligent in collecting evidence; there is evidence I don't have but should have. Furthermore, suppose that if I read the article, I would have acquired strong evidence against (G). Thus, while (G) is supported by the evidence I do have, I should have had additional evidence and I shouldn't believe (G). So, (O2) has the wrong result in this sort of case. It implies that I ought to follow the limited evidence I do have rather than the larger body of evidence I have negligently avoided.[40]

The details of examples such as this one are quite important. There are possible examples in which a person has some good reasons to believe something, but knows that there are available strong considerations to the contrary although he chooses not to make himself familiar with that evidence. Suppose, that the magazine headline in our example is "Gingko Shown to be Ineffective." Upon seeing this headline, I immediately stop reading since I don't want evidence like that. I continue to believe (G) on the basis of my old reasons. Surely my conduct in this case is reprehensible. But just as surely, I have acquired evidence against (G) and my belief losses some considerable support the moment I see the title of the article. It gives me good reason to think that there are strong objections to (G), even if I'm not yet in a position to say in any detail what they are. That significantly alters the evidential status of the proposition for me. Given the credibility of the source and the nature of the article title, most likely my overall evidence no longer supports my belief. I no longer ought to believe (G).

Variations on the example are possible. The article might be in a publication whose reliability is entirely unknown to me or in a magazine I know to be thoroughly disreputable. Other variations on the example concern the wording of the title. It could be something neutral such as "Some New Information on Gingko." Again I ignore the article. Again, assume that it

---

[40] Keith DeRose presented an objection along these lines when he commented on a version of this paper presented at the Rutgers Epistemology Conference.

contains the same devastating objections. In this case, the mere awareness of the existence of the article has much less evidential force. Nevertheless, you might plausibly think that in these cases, or at least some of them, I am terribly negligent in ignoring the article. I shouldn't do that. Had I not done it, I wouldn't have maintained my belief in (G). So, I ought not have that belief.

The statement of the objection seems to me to state the heart of the reply. The name and source of the message do not provide me with much, if any, reason to stop believing (G).[41] Until I've read the article, it would be bizarre for me to stop believing what's supported by my old evidence. I don't have any reason to. Suppose that instead of negligently ignoring the article, I'm busy and I simply set it aside until I can give it proper attention. There's nothing negligent about this behavior. What should I believe in the meantime? Should I think, "Well, somebody has written something about gingko, so I'd better stop believing that it is effective"? I think that the answer is obvious: I shouldn't change for that reason. And this fact doesn't change whether my motives for setting aside the article are laudatory or reprehensible. In every case, the right answer depends upon what evidence I already have, including the evidence about the possible existence of these objections.

No matter what the answers are to questions about how I ought to conduct my inquiry, where I ought to look for evidence, and so on, there always remain the questions, "What should I believe in the meantime?" "What should I believe until I have a chance (or the courage) to look at that new evidence?" It's that natural and central question to which evidentialism provides a good answer.[42]

One might think that when one should look at additional evidence, one should always suspend judgment in the meantime. But that is a clear mistake, as the example above illustrates. You should follow your current evidence. If you do get more evidence, then you should follow the combined evidence you have at that time.

D2. Duty to Gather Evidence

Even if it is true that in the examples just discussed I am believing as I ought given the limited evidence I have, it is plausible to think that there are other epistemic requirements that I ought to fulfill. In particular, in some versions of these cases, I ought to obtain the additional evidence that is available to me. More generally, in a wide variety of cases, it might be thought, one

---

41 One might argue that we all know that magazines usually run articles like this only when they are negative. If that's the case, then seeing the seemingly neutral title does provide me with some evidence against (G). The case is then similar to the first version in which the title is explicitly negative.

42 Richard Foley takes a similar approach to epistemic questions. See especially Chapter 1 of *Working Without a Net* (Oxford: Oxford University Press, 1993).

ought to gather additional evidence about the propositions one considers. It's not enough simply to follow the evidence one actually has. Thus, evidentialism overlooks important epistemic oughts. In a recent paper, "The Epistemic Duty to Seek More Evidence," Richard J. Hall and Charles R. Johnson argue that you have an epistemic duty to seek more evidence about every proposition about which you are not certain.[43] While many epistemologists would reject the demanding condition Hall and Johnson describe, it may seem obvious that in many cases you ought to seek evidence regarding propositions about which you will form beliefs. However, the evidentialist view defended here apparently conflicts with this seemingly obvious truth. If the fundamental epistemic goal is just to have reasonable beliefs, then nothing about evidence gathering techniques or the like follows as a means to that goal.[44] This goal just has implications concerning what attitudes we ought to take given the evidence we have.

By seeking out new evidence concerning some important proposition and then believing what that evidence supports, I don't do a better job of achieving the goal of believing reasonably. I achieve that goal at any moment by believing what is then supported by my evidence. It's surely true that there are times when one would be best off finding new evidence. But this always turns on what options one has, what one cares about, and other non-epistemic factors. As I see it, these are prudential or moral matters, not strictly epistemic matters.

A familiar distinction may help make the evidentialist account seem more attractive. This is the distinction between short-term and long-term goals and the related distinction between synchronic and diachronic rationality. The former concerns questions of rationality at a given moment while the latter concerns rationality over time. Evidentialism is best seen as a theory about synchronic rationality. It holds that the epistemically rational thing to do at any moment is to follow the evidence you have at that moment. It doesn't address questions of how to conduct inquiry over periods of time. Thus, it does not address questions about how to gather evidence, when one ought to seek additional evidence, and so on. In my view these diachronic questions are moral or prudential questions rather than epistemic questions. You should gather more evidence concerning a proposition only when having a true belief about the subject matter of the proposition makes a moral or prudential difference and gathering more evidence is likely to improve your chances of

---

[43] See Richard J. Hall and Charles R. Johnson, "The Epistemic Duty to Seek More Evidence," *American Philosophical Quarterly* 35 (1998): 129–40. Hall and Johnson actually defend the thesis that this remarkable duty follows from the assumption that our goal is to believe all and only true propositions. They do not defend the claim that this is our goal.

[44] Similarly, if the goal is just to now have true beliefs about the propositions about which we are about to form attitudes, then there's no need to gather new evidence.

getting it right. Of course, whether you ought to gather such evidence also depends upon what other things you could do with your time. Epistemological considerations simply don't resolve such matters.

There are cases in which one can spend one's time gathering evidence about propositions concerning inconsequential and trivial propositions or about more weighty matters. Evidentialism provides no guidance about what to do. As I see it, this is not a weakness of evidentialism, since such choices are not to be made on epistemic grounds. What topics you ought to investigate depend upon what topics are of interest to you, what investigations can help you to make your own life or the lives of others better, and other such matters. Evidentialism is silent on those moral and prudential issues, and I don't see why it should address them.

D3. *Being Rational By Avoiding Evidence*

The evidentialist principles stated here imply that one can get oneself into a highly rational state by ridding oneself of as much evidence as one can and then suspending judgment about virtually everything that comes to mind. If a person finds a drug or a machine that can erase memories from his brain and arranges to be immersed in a sensory depravation tank, he'll have very little evidence regarding anything. By believing very little, he'll then be highly rational according to evidentialist standards. Yet, this may seem just the opposite of being rational.[45]

Evidentialists are committed to the view that the person just described would have little evidence regarding anything and would be rational to suspend judgment about nearly everything. But that seems to be exactly the right conclusion. Once the person has lost his evidence, he has no reasons to believe much, and he'd be unreasonable if he did believe things that would have been well justified for him had he been in more normal circumstances. But he's not in normal circumstances and evidentialism concerns itself with assessing what he should believe in the circumstances he is in, not with what he should believe if he were in different circumstances.

Critics might think that he shouldn't get himself in a condition of evidential deprivation in the first place. Evidentialists can agree, though not on evidentialist grounds. By putting himself in those conditions, he's made himself unable to act on behalf of others who may rely on him. Perhaps he morally ought to avoid rendering himself useless in this way. Furthermore, by putting himself in a condition of evidential deprivation, he deprives himself of all the pleasure that comes from experiencing the world. Unless his prior circumstances and future prospects are extremely grim, it is likely to be imprudent to put himself in that situation. Evidentialists can agree with these critical evaluations of this behavior, but these are not epistemic evalua-

[45] This objection was suggested to me by Timothy Williamson.

tions. They are moral or prudential evaluations of behavior related to the formation of beliefs.

I conclude that objections such as the ones considered here do not undermine the evidentialist view of what we epistemically ought to do. Evidentialism says that when adopting a doxastic attitude toward a proposition, a person epistemically ought to adopt that attitude that is supported by the evidence the person has at that time. How the person came to have that evidence, whether by conscientious inquiry or by avoiding potentially troublesome information, is irrelevant to this epistemic fact. Similarly, how the person ought to proceed in future inquiry is also irrelevant. Evaluations of this behavior can be made, of course, but these evaluations are of a different nature than those made on evidentialist grounds. The evidentialist evaluation assesses whether the attitude formed in the circumstances one is actually in is the attitude one ought to have formed in those circumstances.

III. Epistemology and Ethics

There are several questions about the relation of epistemic duties (or oughts) to ethical and other duties. For example, there is in the literature considerable discussion of the idea that there's some important connection between what one is epistemically justified in believing and what one is morally justified in believing. On one side, there is the line of thought suggested by Clifford, according to which believing without epistemic justification is always morally wrong. To defend this thesis, you'd have to argue that such believing is always voluntary or else that involuntary behavior can be morally wrong. I'm not inclined to argue for either of these theses. Other questions concern whether epistemic duties are a "special case" of moral duties and whether there can be cases in which there is a conflict between these different kinds of duties.[46] I won't take up these topics here. I will discuss one puzzling issue in this area.

The question I do want to address can best be raised by attending to a remark Hall and Johnson make in defense of their claim that we have an epistemic duty to seek more evidence for all uncertain propositions. They say that this duty "is probably a pretty weak duty; most moral and prudential duties would trump it in cases of conflict."[47] They quote Chisholm as holding that moral duties always trump epistemic duties. What I want to discuss is whether any sense can be made of the claim that moral duties trump, or do not trump, epistemic duties.

There is no problem with the idea that duties of the same kind can be weighed in importance. For example, a person might have a moral duty to

---

[46] For discussion of the latter question, see Eugene Mills, "The Unity of Justification," *Philosophy and Phenomenological Research* 58 (1998): 27–50. For discussion of the former question, see Susan Haack, "The Ethics of Belief Reconsidered."

[47] "The Epistemic Duty to Seek More Evidence," p. 136.

his students to show up for class and a moral duty to his sick child to stay home and care for her. And one might say that one of these duties outweighs, or trumps, the other. The idea here, I take it, is that there is some scale of moral value and that fulfilling one duty contributes more to that value than does fulfilling the other.

What is far less clear is that there is such a thing as *just plain ought*, as opposed to the various kinds of oughts philosophers have succeeded in distinguishing.[48] To understand the issue, it's important to distinguish two very different points. Suppose you accept the view that it is morally wrong to believe on the basis of insufficient evidence. (Set aside worries about voluntarism.) I take it that at most you would think that the fact that a particular action is a case of believing on insufficient evidence counts towards its being wrong. As some would put it, it is *prima facie* wrong. This leaves open the possibility that this case of believing on insufficient evidence has other features that count in its favor, morally speaking. It might be a case of being supportive of one's family or disregarding some excusable flaws in others. It might have much to be said for it. I have no trouble at all understanding the idea that we can in principle compare the negative moral value of failing to follow one's evidence to the positive moral value of the other features of the case. But here we are comparing two moral factors to reach an overall moral assessment.

I take it that when people say things such as "Moral oughts trump epistemic oughts" they are not saying that the moral weight of epistemic oughts is less than the moral weight of other moral considerations. I believe that what they are saying is that there is some sort of generic ought that somehow encompasses moral considerations, epistemic considerations, and perhaps others, and then weighs them against one another to come up with an overall assessment. This is not any particular kind of ought. It is *just plain ought*. I take Hall and Johnson to be suggesting that when you epistemically ought to gather more evidence and you morally ought to do something else, the moral ought "wins" and you just plain ought to do that other thing. It's this that I just don't understand. Of course, by this I mean to suggest that no one else understands it either. It makes no sense.

I know of no way to establish that the notion of *just plain ought* makes no sense. But I can give some partial defense to my view by noting that it would be a mistake to assume that there must be such a thing as "just plain ought" simply because there are various kinds of oughts. There needn't be anything meaningful about their combination. To see why, consider an analogy. Suppose a child has two dolls. One is short and squat. The other is tall and thin. The child asks you which doll is "bigger". You are faced with a

[48] I learned this way of formulating this question from an unpublished paper by Owen McLeod.

problem. One doll is taller. The other is wider and, let's assume, more voluminous. The child could be asking which is taller, which is wider, or which is more voluminous. Each of those measures could be expressed by "bigger". We needn't assume that there's any meaningful measure that combines these three, yielding "bigger all things considered." Or, more precisely, we could arbitrarily combine the various measures in any way we like. But it would be a mistake to pick any such measure as what really counts as "overall bigger." It would make no sense to ask whether height trumps width or volume trumps height in assessing overall size. We've disambiguated the word "bigger". There's no putting the senses back together, at least no way that isn't simply arbitrary. There's no such thing as "just plain bigger." There are simply the various bigger-than relations.

Consider another example. Suppose we are considering the wealth and strength of two people, A and B. A is stronger, B is wealthier. We ought not assume that there's some proper way of combining these two measures into one, which we might call "strealth". We can of course combine measures of strength and wealth in a limitless number of ways. It might be that some combination would in fact best suit our purposes. But the mere fact that there are two measures, wealth and strength, does not guarantee that there is any proper way to combine them to form strealth. The existence of numerical measures of wealth and strength ought not deceive us here. We could measure one's wealth in rubles and one's strength in terms of the number of grams one can bench press. We could then add those numbers together to get a proposed measure of strealth. But we also could add one's wealth in dollars to the number of pounds one can bench press. That will give a different number. And there are a limitless number of other such sums. Each will yield potential strealth rankings. The different possibilities will have the effect of counting one measure more heavily than the other. If someone were to object to one of these strealth measures on the grounds that it conflicts with the fact that wealth trumps strength in measures of strealth, I think we should say that we have no idea what that person is talking about. We have no independent concept of strealth for which there are correct weightings of the two components.

I think that the same is true of the attempt to compare various kinds of oughts to determine what one just plain ought to do. Suppose that one belief is prudent for me—it will maximize my well being or what I care about or what is important to me—but it is not a belief I epistemically ought to have since I lack evidence for it. We can understand the idea that forming an unjustified belief might be imprudent, since it might foster a bad habit. So, we can imagine weighing the short-term prudential gain against the long term prudential cost. But I can see no values to which we could be appealing when we ask whether the overall prudential benefit trumps the epistemic cost.

For each "ought" there is an associated value. We ought, in the relevant sense, to do the thing that maximizes that value, or perhaps something that does well enough in achieving that kind of value. For example we morally ought to do what maximizes, or produces enough, moral value. We prudentially ought to do what maximizes, or produces enough, prudential value. If there is such a thing as "just plain ought" then there is a value associated with it. The thing we just plain ought to do is the thing that comes out highest, or high enough, according to that measure. It's far from clear what that value would be—it isn't to be identified with any of the more determinate kinds of value and there seems to be no uniquely correct (or range of correct) ways to combine moral, practical, epistemic, and other values. We've disambiguated "ought" and we can't put the various senses back together again. There is no meaningful question about whether epistemic oughts "trump" or are trumped by other oughts.

Stewart Cohen has presented the following objection to this view.[49] Suppose a child comes to her parents with a problem. She is contemplating some action and she's figured out that it is the prudent thing to do but not the moral thing to do. In other words, she realizes that she prudentially ought to do it, but she morally ought not do it. She asks, "What ought I do?" She is clearly looking for a way to resolve the apparent conflict. The view I've defended requires the parents to say that there is no meaningful question here. More precisely, when she uses the word "ought" in her question, she is either using it in one of the disambiguated senses, in which case she already knows the answer to her question, or else she is using it in the sense of "just plain ought", in which case she's not asking a meaningful question. In the latter case, the question is like that of child who wants to know which doll is really bigger.[50]

It is difficult to resist the inclination to want to give the child advice about how to conduct herself. And there is advice that one might give. The question invites further reflection on whether the moral evaluation is really correct, given that the act is harmful to her. The question also invites further reflection on whether the prudential calculation is really correct, given that immoral behavior is also often imprudent. And in the particular case of a child asking a parent such a question, there is the matter of whether the child will have parent's approval whichever way she acts. There is, finally, the question of what the parent predicts he or she would do in similar circumstances. Thus, there are several meaningful questions that one might ask in the sort of case Cohen envisions. However, as I see it, none of them amounts to the question of what one just plain ought to do.

---

[49] In conversation.

[50] Another possibility is that she is using "ought" in yet another more determinate sense. This option seems irrelevant to the present example.

IV. Conclusion

I've argued in this paper that deontological judgments about beliefs do not imply doxastic voluntarism. I've also argued that we do have a form of voluntary control over a substantial number of beliefs, but that this fact is of no great significance for epistemology. Epistemic evaluations do not depend upon whether we have or exercise this voluntary control. I've also argued, or asserted, that what we epistemically ought to do is follow our evidence: when we have a doxastic attitude toward a proposition, it ought to be the attitude that is supported by the evidence we have. I haven't attempted to say anything here about what evidential support is or what it is to have something as evidence. Those are topics for another occasion. I have claimed that there is nothing more that a person epistemically ought to do than to follow her evidence in this way. In particular, activities such as gathering additional evidence for propositions about which one is uncertain are not among the things one epistemically ought to do, even if they are on other grounds highly desirable. Finally, I've expressed skepticism about the meaningfulness of questions about whether or not epistemological considerations are outweighed by moral or prudential considerations in figuring out what we ought to do all things considered.

The topics I've addressed by no means exhaust the issues that could properly be raised under the heading "the ethics of belief." But I hope to have made some small contribution to our thinking about the specific topics I have addressed.[51]

[51] In working on this paper I've benefitted greatly from numerous discussions with Earl Conee, John Bennett, Stewart Cohen, Jonathan Vogel, and all the students in my Fall 1998 epistemology seminar. I'm also grateful to John Greco and Matthias Steup for helpful comments. Earlier versions of the paper have been presented at the University of Rochester and at Rutgers University. I'm grateful to the audiences on both occasions and to Keith DeRose for his provocative comments on the latter occasion.