

Chapter Thirteen

Is Justification Internal?

Internalists claim that only what is internal to our minds can be the sort of thing that justifies our beliefs. Externalists reject this requirement. In what sense could a justifier be claimed to be internal to the mind? The question could be answered by saying that what justifies our beliefs are good reasons, and that only mental items – such as sense experiences, intuitions, introspective states, or memories – qualify as good reasons. The internalist constraint might be justified by claiming that (i) justified belief is responsibly held belief and (ii) only internal factors determine whether a belief is held responsibly. To defend externalism, John Greco rejects (ii). He argues that whether a belief is responsibly held is at least in part a function of its etiology (its causal history), and thus partially an external matter. Another defense of internalism is based on the claim that externalism makes the rejection of skepticism too easy. Greco's reply is that, vis-à-vis the task of rejecting skepticism, externalism is better situated since internalism engenders regress problems and thus makes it impossible to advance an effective anti-skeptical argument. In the last section of his essay, Greco argues that all interesting concepts of epistemic evaluation are externalist because they all make etiology relevant.

Richard Feldman defends internalism. Identifying justification with having good reasons, where whether one has good reasons is solely a function of what mental state one is in, Feldman argues that there are at least two interesting epistemic concepts: consistency among one's beliefs, and a belief's being justified by virtue of being supported by good reasons. Feldman accepts that there is an important epistemically evaluative status – well-foundedness – that makes etiology relevant. But, Feldman argues, if having good reasons is part of well-foundedness, and well-foundedness is necessary for knowledge, then having good reasons is necessary for knowledge and therefore qualifies as interesting.

Contemporary Debates in Epistemology, Second Edition. Edited by Matthias Steup, John Turri, and Ernest Sosa.
© 2014 John Wiley & Sons, Inc. Published 2014 by John Wiley & Sons, Inc.

Justification Is Not Internal

John Greco

1 The Internalism–Externalism Debate in Epistemology

When we say that someone knows something we are making a value judgment – we are saying that there is something intellectually good or right about the person's belief, or about the way she believes it, or perhaps about her. We are saying, for example, that her belief is intellectually better than someone else's mere opinion. Notice that we might make this sort of value judgment even if the two persons agree. Suppose that two people agree that the earth is the third planet from the sun. Nevertheless, we might think that one person knows this while the other person merely believes it. If so, we are making a value judgment – we are saying that there is something intellectually better going on in the case of knowledge. Another way to put the point is to say that knowledge is a normative notion. There is something normatively better about the case of knowledge, as opposed to the case of mere opinion, or even the case of true opinion. Finally, saying that someone knows something is not the only sort of value judgment we can make about her belief. Even if we agree that some belief falls short of knowledge, we might nevertheless judge that it is justified, or rational, or reasonable, or responsible. In each such case, we are saying that there is something normatively better about the case in question, as opposed to the case of mere opinion, or even the case of true opinion.

Internalism in epistemology is a thesis about the nature of this sort of normativity. More precisely, it is a thesis about what sorts of factors determine the epistemically normative (or evaluative) status of belief. Internalists claim that the epistemic status of a belief is entirely determined by factors that are relevantly “internal” to the believer's perspective on things. That is, when a person *S* has some belief *b*, whether *b* is justified (or rational, or reasonable, or responsible) for *S* is entirely a function of factors that are relevantly internal to *S*'s perspective. By contrast, the “externalist” in epistemology denies this. The externalist says that the epistemic status of a belief is *not* entirely determined by factors that are internal to the believer's perspective. When internalism and externalism are characterized this way, a number of things become apparent.

First, internalism is a rather strong thesis, in the sense that it says that epistemic status is *entirely* a function of internal factors. By contrast, the denial of internalism is a relatively weak thesis. Externalism in epistemology holds that *some* factors that are relevant to epistemic status are not internal to the believer's perspective. Second, it is apparent that there are a number of kinds of epistemically normative status, corresponding to a number of kinds of epistemic evaluation. As already noted, we can say that a belief is justified, or rational, or reasonable, or intellectually responsible, and these need not mean the same thing. It is possible, then, to be an internalist about some kinds of epistemic status and an externalist about others. Hence there are a variety of internalisms and a corresponding variety of externalisms.

Third, we get different understandings of internalism (and externalism) depending on different ways that we understand “internal to *S*'s perspective.” The most common

way to understand the phrase is as follows. Some factor *F* is internal to *S*'s perspective just in case *S* has some sort of privileged access to whether *F* obtains. For example, a factor *F* is relevantly internal to *S*'s perspective if *S* can know by reflection alone whether *F* obtains.¹ A related, though not equivalent, understanding of "internal to *S*'s perspective" is as follows. Some factor *F* is internal to *S*'s perspective just in case *F* constitutes part of *S*'s mental life.² For example, a person's perceptual experience counts as internal on this understanding, since how things appear perceptually to *S* is part of *S*'s mental life in the relevant sense. Also, any belief or any representation that *S* has about how things are would be internal on this understanding, since a person's beliefs and other representations are also part of her mental life. These two understandings are related because it is plausible to think that one has privileged access to what goes on in one's mental life, and perhaps *only* to what goes on in one's mental life. In that case, the two understandings would amount to the same thing for practical purposes. Internalism would then be the thesis that epistemic status (of some specified sort) is entirely a function of factors that are part of one's mental life, and to which one therefore has privileged access.

Finally, it is apparent that some varieties of internalism are initially more plausible than others. That is, some sorts of epistemic evaluation are obviously externalist on the above understandings. Most importantly, and perhaps most obviously, whether a belief counts as *knowledge* is an external matter, if only because a belief counts as knowledge only if it is true, and whether a belief is true is typically an external matter.

There is another reason why knowledge and many other sorts of epistemic evaluation must be understood as externalist, however. Consider that we can evaluate both persons and their beliefs in two very different ways. Broadly speaking, we can evaluate persons and their beliefs either from an "objective" point of view or from a "subjective" point of view. From the objective point of view, we can ask whether there is a "good fit" between the person's cognitive powers and the world. For example, we can ask whether the person has a sound understanding of the world around her, or whether she has a good memory, or accurate vision. Also from this point of view, we can ask whether a person's methods of investigation are "reliable," in the sense that they are likely to produce accurate results. Notice that when a person gets positive evaluations along these dimensions, her relevant beliefs will as well: her beliefs will be largely true, or objectively probable, or objectively well formed, or reliably formed. On the other hand, there is a second broad category of epistemic evaluation. This sort concerns not whether a belief is objectively well formed, but whether it is subjectively well formed. It asks not about objective fitness, but about subjective appropriateness.

Common sense tells us that these two sorts of evaluation can come apart. For example, suppose that someone learns the history of his country from unreliable testimony. Although the person has every reason to believe the books that he reads and the people that teach him, his understanding of history is in fact the result of systematic lies and other sorts of deception. How should we evaluate this person's beliefs epistemically? By hypothesis, they are not well formed, objectively speaking: they are based on lies and deceptions. Nevertheless, there are clear senses in which the person's beliefs might be subjectively well formed. If the person has been deceived through no fault of his own, we might fairly say that his beliefs are intellectually responsible, or perhaps epistemically rational. We are inclined to say similar things about the victim of a convincing hallucination. Suppose that Descartes believes there

is a fire before him, and that he believes this on the basis of vivid sensory experience embedded in a broad and coherent set of background beliefs. But suppose also that Descartes is the victim of a massive and systematic illusion. The illusion, we can imagine, is undetectable and occurs through no fault of his own. Again, any epistemic evaluation of Descartes's belief will fall into one of two broad categories, or some combination of these. The belief can be evaluated in terms of its objective fit (in which case it fares poorly), or it can be evaluated in terms of its subjective appropriateness (in which case it fares well).

And now to the point: internalism is pretty much a nonstarter with respect to evaluations of the first category. Evaluations from an objective point of view involve factors such as accuracy, reliability, and appropriate causal relations to one's environment. And these are paradigmatically external factors. That is, they are factors that cannot be understood as internal to our cognitive perspective, whether we understand "internal" in terms of privileged access or in terms of what goes on in one's mental life. This is why there are no internalist theories of knowledge. Knowledge, it would seem, requires both objective and subjective factors. Put another way, a belief counts as knowledge only if it is *both* objectively well formed and subjectively appropriate. But since the former sort of status involves external factors, knowledge itself is external. Internalism, therefore, is best understood as a thesis about the second broad category of epistemic evaluation: it is a thesis about what factors determine subjective appropriateness.

We have now arrived at the following understanding of internalism:

- (I) Whether a belief *b* is subjectively appropriate for a person *S* is entirely a matter of factors that are internal to *S*'s perspective.³

Suppose we use the term "epistemic justification" to name the sort of subjective appropriateness that is required for knowledge. A standard form of internalism says that (I) holds with regard to epistemic justification (for example, see Ginet, 1975; Chisholm, 1977; Conee and Feldman, 2001). Alternatively, one might think that (I) holds for other kinds of subjective appropriateness, independently of their connection to knowledge. For example, one might think that rationality is an important and independent normative property, and that whether a belief is rational is entirely a matter of factors that are internal to *S*'s perspective.

I will argue that internalism is false in all of its varieties. More exactly, I will argue that internalism is false in all of its *interesting* varieties – it is false as a thesis about any interesting or important sort of normative epistemic status. Most importantly, internalism is false as a thesis about epistemic justification, or the kind of subjective appropriateness that is required for knowledge. That is bad enough, but in fact the situation for internalism is much worse. Internalism is false as a thesis about any interesting or important sort of epistemic evaluation, and any corresponding sort of epistemic normativity. In section 2, I will look at three considerations that are commonly put forward in favor of internalism as a thesis about epistemic justification, and I will argue that none of these adequately motivates the position. In fact, all three considerations motivate externalism about epistemic justification. In section 3, I will give a general argument against internalism in all its (interesting) varieties.

2 Three Motivations for Internalism

Three considerations are commonly put forward in favor of internalism, where internalism is understood as a thesis about epistemic justification, or the sort of justification that is required for knowledge. I will consider these in turn.

2.1 Epistemic justification as epistemic responsibility

The first consideration begins with an assumption about the nature of epistemic justification. Namely, a belief *b* is epistemically justified for a person *S* just in case *S*'s belief *b* is epistemically responsible. However, the argument continues, epistemic responsibility is entirely a matter of factors that are internal to *S*'s perspective. Therefore, epistemic justification is entirely a matter of factors that are internal to *S*'s perspective.⁴

The essentials of the argument can be stated this way:

1. A belief *b* is epistemically justified for a person *S* just in case *S*'s believing *b* is epistemically responsible.
2. Epistemic responsibility is entirely a matter of factors that are internal to *S*'s perspective.

Therefore,

3. Epistemic justification is entirely a matter of factors that are internal to *S*'s perspective. (1, 2)

Let us grant for the sake of argument that epistemic justification is a matter of epistemic responsibility. In other words, let us grant premise 1 of the argument. Nevertheless, premise 2 of the argument is false. Specifically, it is not true that epistemic responsibility is entirely a matter of factors that are internal to *S*'s perspective. Two sorts of considerations establish this point.

First, the notion of responsibility is closely tied to the notions of blame and praise. For example, judgments concerning whether a person is *morally* responsible with respect to some action or event are often equivalent to judgments about whether the person is morally blameworthy with respect to the action or event. Similarly, judgments concerning whether a person is *epistemically* responsible with respect to some belief *b* are often equivalent to judgments about whether the person is epistemically blameworthy with respect to *b*. And now the point is this: whether a person is epistemically blameworthy for holding some belief is partly a function of the person's prior behavior: if *S*'s reasons for believing *b* are the result of prior negligence, then *S* is not now blameless in believing *b*. An example will illustrate the point.

Example 1. Maria believes that Dean Martin is Italian. She believes this because she seems to remember clearly that it is so, and she presently has no reason for doubting her belief. But suppose also that Maria first came to this belief carelessly and irresponsibly (although she has now forgotten this). Many years ago, she formed her belief on the basis of testimony from her mother, who believes that all good singers are Italian. At the time Maria knew that her mother was an unreliable source in these matters, and she realized that it was not rational to accept her mother's testimony.

Clearly, Maria is not now blameless in believing that Dean Martin is Italian. Again, prior negligence is a factor determining present responsibility, and even if that negligence is not internal to S's perspective. Therefore, premise 2 of the argument above is false.

A second consideration also establishes that premise 2 is false. First, we can make a distinction between (a) *having* good reasons for what one believes, and (b) believing *for* good reasons. Anyone who knows the axioms of arithmetic *has* good reasons for believing a theorem in the system. But unless one puts two and two together, so to speak, one does not believe the theorem in question *for* the right reasons.⁵ And now the point is this: a belief is epistemically praiseworthy only if it is believed for the right reasons. Two examples illustrate the point.

Example 2. A math student knows all the relevant axioms but doesn't see how the axioms support a theorem that must be proven on the exam. Eventually he reasons fallaciously to the theorem, and believes it on the basis of his fallacious reasoning.

Example 3. Charlie is a wishful thinker and believes that he is about to arrive at his destination on time. He has good reasons for believing this, including his memory of train schedules, maps, the correct time at departure and at various stops, and so forth. However, none of these things is behind his belief – he does not believe what he does *because* he has these reasons. Rather, it is his wishful thinking that causes his belief. Accordingly, he would believe that he is about to arrive on time even if he were not.

Clearly, the math student's belief about the theorem is not praiseworthy. Likewise, Charlie is not praiseworthy in believing that he will arrive on time.

The moral to draw from both examples is that "etiology matters" for epistemic responsibility. In other words, whether a belief counts as epistemically responsible depends, in part, on how the belief was formed. Since these beliefs were formed on the basis of bad reasons rather than good reasons, they are not epistemically praiseworthy. In fact, the same moral can be applied to example 1. Prior negligence also figures into the etiology of a belief, and is a factor in determining whether a belief is epistemically responsible. And of course, the etiology of a belief concerns factors that are external to the believer's perspective. Putting all this together, we may conclude that epistemic responsibility is not entirely a matter of factors that are internal to S's perspective. Accordingly, understanding epistemic justification in terms of epistemic responsibility does not motivate internalism about epistemic justification. In fact, it motivates externalism about epistemic justification (see Greco, 1990).

2.2 Like believers have like justification

A second consideration that is sometimes put forward in favor of internalism invokes a strong intuition about epistemic justification. Namely, in many cases it seems that believers who are alike in terms of internal perspective must also be alike in terms of epistemic justification. The point is often illustrated by considering Descartes's victim of an evil deceiver. Suppose that the victim is exactly like you in terms of internal perspective. Even if the victim lacks knowledge, the argument goes, surely his beliefs are as well justified as yours are. If you are justified in believing that there is a table before you, and if the victim's perspective is exactly as yours, then he must be justified in believing that there is a table before him.⁶

The considerations about epistemic responsibility above suffice to counter this line of reasoning, however. The problem is that two believers might be alike internally, and yet different regarding the causal genesis of their beliefs. And once again, etiology matters. Suppose that two persons arrive at the same internal perspective, but that one does so in a way that is epistemically responsible, whereas the other does so on the basis of carelessness, thick-headedness, and stupidity. The two persons will not be alike in epistemic justification, although they share the same internal perspective.

2.3 Replying to skepticism

A third consideration invoked in favor of internalism is that externalism makes an answer to skepticism too easy. Philosophical problems are supposed to be difficult. If the externalist has an easy answer to the problem of skepticism, this argument goes, then that is good reason to think that externalism is false. At the very least, it is good reason to think that the externalist has changed the subject, that he is no longer talking about our traditional notions of justification and knowledge.⁷

How does externalism make an answer to skepticism too easy? The idea is roughly as follows. According to the skeptic, one can know via sense perception only if one knows that sense perception is reliable. Similarly, one can know by inductive reasoning only if one knows that inductive reasoning is reliable. This creates problems for the internalist, because it is hard to see how one can mount a non-circular argument to the desired conclusions about the reliability of one's cognitive powers. There is no such problem for the externalist, however, since the externalist can deny the initial assumption of the skeptical argument. For example, an externalist can insist that sense perception gives rise to knowledge so long as sense perception *is* reliable. There need be no requirement, on an externalist account, that one *know* that one's perception is reliable. What is more, on an externalist account one seemingly *can* know that one's cognitive powers are reliable, and easily so. For example, one can use reliable perception to check up on perception, and then reason from there that perception is reliable. Similarly, one can use reliable induction to check up on induction, and then reason from there that induction is reliable.

In this context Richard Fumerton writes,

All of this will, of course, drive the skeptic crazy. You cannot *use* perception to justify the reliability of perception! ... You cannot *use* induction to justify the reliability of induction! Such attempts to respond to the skeptic's concerns involve blatant, indeed pathetic, circularity. (Fumerton, 1995, p. 177)

The fundamental objection to externalism can be easily summarized. If we understand epistemic concepts as the externalist suggests we do, then there would be no objection in principle to using perception to justify reliance on perception ... and induction to justify reliance on induction. But there is no philosophically interesting concept of justification or knowledge that would allow us to use a kind of reasoning to justify the legitimacy of using that reasoning. Therefore, the externalist has failed to analyze a philosophically interesting concept of justification or knowledge. (Fumerton, 1995, p. 180)

The problem with this argument against externalism is that it is self-defeating. In effect, the argument claims that only internalism can give a satisfying reply to traditional

skeptical concerns. On the contrary, I want to argue, if one concedes internalism then it is *impossible* to give a satisfying reply to traditional skeptical concerns. Specifically, if one concedes that epistemic justification is internalist, then the skeptic has all he needs to construct skeptical arguments that are otherwise sound. Put simply, internalism about epistemic justification guarantees skepticism about epistemic justification.

We may see the point if we consider non-basic knowledge, or knowledge based on evidential grounds.⁸ Presumably, S knows p on the basis of evidence E only if E is a reliable indication that p is true. For example, consider the case where S believes that there is a bird in the tree on the basis of her sensory evidence. Presumably, S *knows* that there is a bird in the tree only if the sensory evidence that S has is indeed a *reliable* indication that there is a bird in the tree. That seems to be something that anyone should concede. But now something important follows from this. Namely, the reliability of S's evidence is one factor that determines whether S's belief has epistemic justification. On the assumption of internalism, then, S knows p on the basis of E only if E's reliability is something that is within S's perspective.

But now what does *that* mean? In what sense could the reliability of one's evidence be within one's perspective? Presumably, the internalist will have to accept something like this: in cases of knowledge, S's belief that "E is a reliable indication that p" is itself epistemically justified. But now how could that be? How could S have epistemic justification for this belief about her evidence? On the assumption of internalism, it is hard to see how she could.

Consider propositions of the form "E is a reliable indication that p." For example, consider the belief that such and such sensory evidence is a reliable indication that there is a bird in the tree. Clearly, this is itself a belief about the world. That is, it is a belief about the character of one's sensory appearances, and about the relationship between those sorts of appearances and real birds and trees. Now this sort of belief will not be knowable a priori. Rather, it is the sort of belief that is known, if at all, on the basis of empirical evidence. And therefore we are threatened with a regress or a circle. That is, if S can know that there is a bird in the tree only if she knows that her evidence for this is reliable, and if she can know that her evidence is reliable only if both (a) she has evidence for *this*, and (b) she has evidence that this new evidence is reliable, then there would seem to be no end to this sort of problem. For presumably S's belief that her new evidence is reliable will require further evidence, and it will now be necessary that S know that *this* evidence is reliable, and so on.

This problem was illustrated by the example of knowing that there is a bird in the tree on the basis of sensory appearances. But it is really a very general problem, which arises in any case where a belief about the world is based on empirical evidence. For in any such case, the belief that one's evidence is reliable will itself be a belief about the world based on empirical evidence, and so we will be off and running. For example, consider the belief that all crows are black, which is based on inductive evidence involving past observations of crows. If internalism is true, then one is justified in believing that all crows are black only if one is justified in believing that one's past observations are a reliable indication of one's present belief about crows. But this belief about one's evidence is itself a belief about the world, and will itself require empirical evidence.

One strategy for avoiding this sort of problem is to attempt an a priori argument to the effect that one's evidence is reliable. That is, one might try to show that one's evidence is reliable, but without using further empirical evidence to do so. That would

stop the regress in its tracks. But this strategy is a dead end. In principle, it would require showing that our sensory evidence *must* be a reliable indication of our perceptual beliefs, and that our inductive evidence *must* be a reliable indication of our inductive beliefs. But neither of these things is true. Rather, it is at most a contingent fact about us and our world, not a necessary fact about our evidence, that sensory appearances are a reliable indication of perceptual beliefs. Likewise, it is at most a contingent fact about us and our world, not a necessary fact about our evidence, that past observations are a reliable indication of unobserved cases.

The line of reasoning set out above is closely analogous to Hume's skeptical reasoning. Just like our internalist, Hume believed that one's empirical evidence gives rise to knowledge only if one knows that one's evidence is reliable. For example, Hume thought one must know that, in general, sensory appearances are a reliable guide to reality. Likewise, he thought one must know that, in general, observed cases are a reliable indication of unobserved cases. But there is no way to know such things, Hume argued, without reasoning in a circle. And so there is no way to know such things at all. The present point is this: if one adopts an internalist account of epistemic justification, then Hume has all the premises he needs to mount his skeptical argument. Put another way, there will be nothing else to challenge in Hume's reasoning. Here is that reasoning, set out more formally:

Skepticism about perception

- 1 All our perceptual beliefs depend for their evidence on (a) sensory appearances, and (b) the assumption (R) that sensory appearances are a reliable indication of how things are in the world.
- 2 But (R) is itself an assumption about how things are in the world, and so ultimately depends for its evidence on perceptual beliefs involving sensory appearances.
- 3 Therefore, assumption (R) depends for its evidence on (R). (1, 2)
- 4 Circular reasoning cannot give rise to knowledge.
- 5 Therefore, (R) is not known. (3, 4)
- 6 All our perceptual beliefs depend for their evidence on an assumption that is not known. (1, 5)
- 7 Beliefs that depend for their evidence on an unknown assumption are themselves not known.
- 8 Therefore, no one has perceptual knowledge. (6, 7)

Skepticism about induction

- 1 All our inductive beliefs depend for their evidence on (a) past and/or present observations, and (b) the assumption (R') that observed cases are a reliable indication of unobserved cases.
- 2 But (R') is itself an assumption about something unobserved, and so ultimately depends for its evidence on induction from past and/or present observations.
- 3 Therefore, belief (R') depends for its evidence on (R'). (1, 2)
- 4 Circular reasoning cannot give rise to knowledge.
- 5 Therefore, (R') is not known. (3, 4)
- 6 All our beliefs about unobserved matters of fact depend for their evidence on an assumption that is not known. (1, 5)

- 7 Beliefs that depend for their evidence on an unknown assumption are themselves not known.
- 8 Therefore, no one knows anything about unobserved matters of fact. (6, 7)

What is wrong with Hume's arguments? In each case, the independent premises of the argument are 1, 2, 4, and 7. Also in each case, 2, 4, and 7 seem uncontroversial. That leaves premise 1 as the only thing left to challenge. But if internalism is true, then in each case premise 1 is true.

Once again, a closer look at a motivation for internalism ends up providing a motivation for externalism. Internalism, we have seen, makes it impossible to reply to Hume's skeptical arguments. On the contrary, we can hope to avoid Hume's skeptical conclusions only by adopting externalism.⁹

3 The General Argument against Internalism

The arguments in section 2 are directed against internalism about epistemic justification, or the sort of subjective justification that is required for knowledge. We have noted, however, that it is possible to be an internalist about other kinds of normative epistemic status. More specifically, one might think that there are other sorts of subjective appropriateness, which are independent of knowledge and epistemic justification, but which nevertheless correspond to interesting and important kinds of epistemic evaluation. Perhaps this sort of normativity is at issue when we evaluate the beliefs of Descartes's demon victim. There is no question about whether the victim has knowledge, and we might even agree that he lacks epistemic justification. Still, we might insist, there is *some* sense in which the victim's beliefs are intellectually respectable, or up to par, or at least not objectionable. In this section I will argue that there is no important or interesting normative property that is also internalist. To be clear, I do not mean to deny that there are important epistemic properties that are independent of knowledge and justification. I think there are. My point is rather that no such property is internalist.¹⁰

The argument begins by recalling the two broad kinds of evaluation noted above. We said that, broadly speaking, we can evaluate persons and their beliefs either from an "objective" point of view or from a "subjective" point of view. The objective point of view concerns (roughly) whether a belief has good objective fit with the world. From this point of view, we ask such questions as whether a belief is accurate, or reliably formed, or appropriately causally related to the facts. The subjective point of view concerns (roughly) what is subjectively appropriate to believe. From this point of view, we ask such questions as whether a belief is subjectively plausible, or responsibly formed, or well motivated. The argument then proceeds as follows. First, evaluations from the objective point of view are obviously externalist. Considerations concerning accuracy, reliability, and causal relations involve factors that are paradigmatically externalist. But, second, evaluations from the subjective point of view are also externalist. For example, the considerations in section 2 show that epistemic responsibility, no less than reliability, is a function of etiology. But the etiology of a belief is an external matter – it concerns such things as the history of the belief and the reasons why it is held, and these are things that are typically external to one's perspective.

And now for the final premise of the argument: there is no interesting or important kind of epistemic evaluation that does not concern either objective fit or subjective appropriateness. Therefore, no interesting or important kind of epistemic evaluation, and no corresponding sort of epistemic normativity, is internalist.

A corollary of this argument is that all interesting kinds of epistemic normativity depend on factors related to accuracy and/or etiology. Of course, one can *stipulate* a kind of evaluation that abstracts away from these entirely. For example, we can stipulate that Mary's belief about Dean Martin is "weakly blameless" in the following sense: S's belief *b* is weakly blameless just in case S is no more blameworthy at the moment for believing *b* than she was a moment before. The present point is that this sort of normativity will not be interesting.

One way to see why is to look at the purpose of epistemic evaluation. It has often been noted that knowledge is a social product with a practical value. We are social, highly interdependent, information-using, information-sharing beings. As such, it is essential to our form of life that we are able to identify good information and good sources of information. In this context, it is not surprising that we make evaluations concerning how beliefs are formed, their history in relation to other beliefs, why they are believed, and so on. In other words, it is not surprising that we make evaluations concerning whether beliefs are reliably and responsibly formed. But evaluations of these sorts involve considerations about accuracy and etiology. And, therefore, evaluations of these sorts are externalist evaluations.

This context also shows why judgments that abstract entirely away from external factors will be uninteresting. For example, why should we care that Mary is no more blameworthy for her belief at the moment than she was the moment before? We care about whether Mary is, in general, a responsible and reliable cognitive agent. We also care about whether, in this instance, Mary arrived at her belief in a reliable and responsible way. We also care, of course, about whether Mary's belief is true. These are important considerations about Mary and about her belief – considerations that are important from the point of view of information-using, information-sharing beings such as ourselves. On the other hand, "time-slice" evaluations that abstract away entirely from the formation of beliefs, their relation to the world, and the character of believers will not be very important. Of course, we often want to abstract away from *some* external factors – we often want to abstract away from some or others. The point here is that we never want to abstract away from *all of them at once*. In other words, we have no interest in epistemic evaluations that are (entirely) internalist.

Consider the analogy to moral evaluation. We care about which people are good and which actions are right. That is, we care whether, in general, a person is a reliable and responsible moral agent. And we care about whether, in a particular instance, a person acted in a responsible and reliable way. What we don't care about is artificial, time-slice evaluations such as that S is not more blameworthy at the moment for bringing about some state of affairs than she was the moment before. Neither do we care whether some action *A* is right relative to S's own moral norms, in abstraction from questions about how S did *A*, or why S did *A*, or whether S's norms are themselves any good. Of course, we often want to abstract away from *some* external considerations – we want to abstract away from some or others. The point is that we never want to abstract away from *all of them at once*. In other words, we have no interest in moral evaluations that are (entirely) internalist.

These last points can be illustrated by applying them to a particular version of internalism. A number of philosophers have wedded internalism to evidentialism (for example, see BonJour, 1985; Conee and Feldman, 2001). The main idea behind evidentialism is that a belief has positive epistemic status (of some sort or another) if and only if it is appropriately related to good evidence. Put another way, S's belief *b* has positive epistemic status just in case *b* "fits" S's evidence. The internalist adds to this that notions such as "good evidence" and "fit" are to be understood along internalist lines (for example, see Conee and Feldman, 2001, 1985).

Consider a case where S has a belief *b* and evidence *E*. For example, let *E* be a set of observations together with relevant background beliefs, and let *b* be the belief that all crows are black. There are many dimensions along which S's belief can be evaluated. For example, we can ask (a) whether *E* is true, (b) whether *E* is objectively probable, (c) whether *E* was reliably formed, (d) whether *E* was responsibly formed, (e) whether *E* leaves important information out, (f) whether *E* is a reliable indication that *b* is true, (g) whether *b* is objectively probable on *E*, and (h) whether S believes *b* because S believes *E*. All of these are external matters, involving factors that are neither part of S's mental life nor something to which S has privileged access in the typical case. Now certainly, in some situations we want to abstract from some of these matters. For example, we might be interested to know whether S's evidence is responsibly formed, whether or not S's evidence is true. Alternatively, we might be interested to know whether S's evidence makes her belief probable, whether or not her evidence was responsibly formed. But is there any situation in which we are interested to abstract away from *all* these factors at once? In other words, is there any situation in which we are interested to abstract away from *all* external factors? It is hard to imagine that there is.

Let us consider what is perhaps the most plausible possibility along these lines. We are sometimes interested to know, it might be suggested, whether S's belief is justified in the following sense: S believes *b* and S has *E*, and believing *b* on the basis of *E* would be licensed by the norms of evidence that S accepts. Remember, however, that we are supposed to be abstracting away from all external considerations. There is no question, therefore, whether S's norms of evidence are in fact reliable, or whether *E* is itself probable, or whether *E* was arrived at through prior negligence, or whether S believes *b* because S believes *E*, and so on. Abstracting away from all of this, why would we be interested to know whether *b* is licensed by norms of evidence that S accepts? Why would this be an important evaluation to make? It would be analogous to asking whether S's action *A* is licensed by the *moral* norms that S accepts, but independently of any questions about the adequacy of S's moral norms, or prior negligence by S, or the probable or actual consequences of *A*, or S's motives in performing *A*. As in the epistemic case, it is hard to imagine a situation in which that sort of moral evaluation would be interesting or important. Both moral and epistemic evaluations, we may conclude, are more closely tied to the world than that. They concern not just what is internal to one's perspective, but how that perspective is related to things outside it.

One apparent drawback of this argument is that it is difficult to fight over what is "important" or "interesting." However, in my mind the dispute between internalists and externalists comes down largely to just this issue. Accordingly, I have tried to put the ball in the internalist court by (a) focusing the argument right there and (b) making it explicit that internalist evaluations abstract away from *all* externalist considerations.

When we focus the dispute in this way, I think many will agree that the sorts of evaluations left over – the sorts that count as internalist – are not very interesting. Not in life, because they do not serve the purposes of information-using, information-sharing cognitive agents. Not in philosophy, because an internalist reply to traditional skeptical arguments is impossible, and so cannot serve that philosophical purpose.

4 Conclusion

We may now take stock of the arguments presented against internalism and in favor of externalism. In section 2, we looked at three considerations that are commonly put forward as motivations for internalism about epistemic justification, or the kind of justification required for knowledge. In each case, we saw that the consideration in question failed to motivate internalism. In fact, each consideration motivated externalism about epistemic justification. In section 3, we considered a general argument against internalism in all its interesting varieties. The argument was that all interesting epistemic evaluations are made from the objective point of view or the subjective point of view – they concern questions about objective fit, or subjective appropriateness, or both. But all such evaluations involve considerations about the accuracy of beliefs and/or their etiology, and these are paradigmatically externalist factors. The conclusion is that there are no interesting internalist evaluations. Put another way, all interesting epistemic evaluations are externalist evaluations.

Acknowledgments

Thanks to Richard Feldman, Stephen Grimm, and Ernest Sosa for helpful comments and discussion.

Notes

- 1 This is what Alston calls “access internalism” (see Alston, 1985, pp. 57–89; 1986, pp. 179–221). See also Ginet (1975), Chisholm (1977), and Goldman (1980, pp. 27–51).
- 2 This is Earl Conee’s and Richard Feldman’s understanding (see Conee and Feldman, 2001, p. 233).
- 3 If we understand “internal to S’s perspective” in terms of privileged access, we get:

(I-PA) Whether *b* is subjectively appropriate for *S* is entirely a matter of factors to which *S* has a privileged epistemic access.

If we understand “internal to S’s perspective” in terms of what goes on in S’s mental life, we get:

(I-M) Whether *b* is subjectively appropriate for *S* is entirely a matter of factors that constitute part of S’s mental life.

In the remainder of the essay I will ignore this distinction, since it is not important to the arguments that follow.

- 4 For arguments along these lines see Ginet (1975) and BonJour (1985).
- 5 For an extended discussion of this distinction and its importance, see Audi (1993, esp. chapter 7).
- 6 For arguments along this line see Foley (1984) and Luper-Foy (1988, p. 361).

- 7 See Fumerton (1995, esp. chapter 6). BonJour (1985, pp. 36–37 and 56–57) hints at this sort of objection. See also Cohen (2002).
- 8 Many internalists believe that all knowledge requires evidential grounds, in which case the following considerations apply to knowledge in general.
- 9 For further discussion along these lines see Greco (1999, 2000). An alternative response to Hume might challenge premise 2 of each argument, claiming that the assumption in question is known because it is reliably formed or because it meets some other externalist criterion for justification and knowledge. It is hard to see how the internalist could pursue an analogous strategy, however. For further discussion of this alternative response to Hume, see Greco (2002, esp. section V).
- 10 Another possibility is that there is some interesting status that is a component of epistemic justification. The arguments below will count against this possibility as well.

References

- Alston, W. (1985) Concepts of epistemic justification. *The Monist* 68: 57–89.
- Alston, W. (1986) Internalism and externalism in epistemology. *Philosophical Topics* 14: 179–221.
- Alston, W. (1989) *Epistemic Justification*. Ithaca, NY: Cornell University Press.
- Audi, R. (1993) *The Structure of Justification*. Cambridge: Cambridge University Press.
- BonJour, L. (1985) *The Structure of Empirical Knowledge*. Cambridge, MA: Harvard University Press.
- Chisholm, R. (1977) *Theory of Knowledge*, 2nd edn. Englewood Cliffs, NJ: Prentice Hall.
- Cohen, S. (2002) Basic knowledge and the problem of easy knowledge. *Philosophy and Phenomenological Research* 65: 309–329.
- Conee, E. and Feldman, R. (2001) Internalism defended. In *Epistemology: Internalism and Externalism*, ed. H. Kornblith. Oxford: Blackwell.
- Conee, E. and Feldman, R. (1985) Evidentialism. *Philosophical Studies* 48: 15–34.
- Foley, R. (1984) Epistemic luck and the purely epistemic. *American Philosophical Quarterly* 21: 113–114.
- Fumerton, R. (1995) *Metaepistemology and Skepticism*. Lanham, MD: Rowman & Littlefield.
- Ginet, C. (1975) *Knowledge, Perception and Memory*. Dordrecht: Reidel.
- Goldman, A. (1980) The internalist conception of justification. *Midwest Studies in Philosophy* 5: 27–51.
- Greco, J. (1990) Internalism and epistemically responsible belief. *Synthese* 85: 245–277.
- Greco, J. (1999) Agent reliabilism. In *Philosophical Perspectives* 13, ed. J. Tomberlin. Atascadero, CA: Ridgeview Press.
- Greco, J. (2000) *Putting Skeptics in Their Place*. Cambridge: Cambridge University Press.
- Greco, J. (2002) How to Reid Moore. *Philosophical Quarterly* 52: 544–563. Reprinted in *The Philosophy of Thomas Reid* (2003), ed. J. Haldane and S. Read. Oxford: Blackwell.
- Luper-Foy, S. (1988) The knower, inside and out. *Synthese* 74: 361.

Justification Is Internal

Richard Feldman

Internalism in epistemology is a view about what sorts of things determine or settle epistemic facts. As its name suggests, it holds that “internal” things determine epistemic

facts. To say that the internal things “determine” the epistemic facts is to say that if two things are alike with respect to internal factors, then they must also be alike in the relevant epistemic ways. But we are not clear about just what internalism amounts to until we are clear about what counts as an internal thing and which epistemic facts are in question. In this essay I will say a little about what counts as an internal thing and considerably more about which epistemic facts are, according to internalism, determined by them. I will argue that internal factors do determine some epistemic facts, and thus my conclusion will be in stark contrast to the thesis John Greco defends in his essay, “Justification Is Not Internal.”

I Background

To understand and assess the current debate about internalism and externalism in epistemology, it is useful to look briefly at the way in which the debate arose. A plausible starting point for this examination is the traditional analysis of knowledge and its rivals.

A The traditional analysis of knowledge

In Plato’s dialogue *Meno* Socrates says, “That there is a difference between right opinion and knowledge is not at all a conjecture with me but something I would particularly assert that I know.” Lucky guesses make vivid the difference to which Socrates refers. If I am a contestant on a game show, I might have the belief that the valuable prize is behind door 1 rather than door 2 or door 3. Suppose my belief turns out to be true. Still, it is clear that I did not know that the prize was behind door 1 until the door was opened and I saw it. This prompts the question, “What is the difference between mere true belief and knowledge?”

Many philosophers thought that knowledge required true belief plus good reasons or good evidence. Thus, Norman Malcolm (1963) wrote, “Whether we should say that you knew, depends in part on whether you had grounds for your assertion and on the strength of those grounds.” Roderick Chisholm (1957, p. 16) wrote that “‘S knows that *h* is true’ means: (i) S accepts *h*; (ii) S has adequate evidence for *h*; and (iii) *h* is true.” These remarks are entirely typical.

The two quotations just presented reveal a generally shared outlook on what is required for knowledge in addition to true belief. Philosophers put the point in a variety of ways. Some say that knowledge requires adequate evidence, others say that it requires good reasons, and others say that it requires strong grounds. And, more commonly, philosophers have expressed this idea by saying that a true belief must be *justified* in order to be knowledge. Thus, a standard way to formulate the traditional analysis of knowledge is by saying that knowledge is justified true belief. This, then, is our starting point. I will sometimes refer to the traditional analysis as “the good reasons” analysis.

There are difficult issues to resolve if the traditional analysis is to be spelled out in detail. For one thing, one must say how much evidence is adequate (or how good the reasons must be or how strong the grounds must be). This detail will not affect the discussion that follows. Another issue has to do with what counts as reasons (or evidence or grounds). This will matter in what follows. It will be discussed briefly in section C, after some rivals to the traditional analysis are introduced.

B Rival accounts of knowledge

Some philosophers have proposed analyses of knowledge that contrast sharply with the “good reasons” type analyses just described. According to one of these alternatives, knowledge requires true belief plus an appropriate type of causal connection to the fact that the belief is about. The details of the causal theory need not concern us here. In the words of one of its leading proponents, a central fact about it is that it “flies in the face of a well established tradition in epistemology, the view that epistemological questions are questions of logic or justification, not causal or genetic questions” (Goldman, 1967, p. 372). David Armstrong (1973, p. 166) defended a similar view. He said that the difference between knowledge and true belief was that knowledge involved “a *law-like connection* between the state of affairs [of a subject’s believing that p] and the state of affairs that makes ‘p’ true such that, given the state of affairs of [the subject’s believing that p], it must be the case that p.” The idea, then, is that whether you have knowledge depends not on what your reasons are, but rather on what the cause of your belief is. People know about the world because there is a causal connection between states in the world and their beliefs. For example, a person can know that it is cold out because the person’s belief that it is cold out is caused by its being cold out.

Another theory that differs significantly from the traditional good reasons view is reliabilism. According to a simple version of reliabilism, a person has knowledge when the person’s true belief is caused by a method of belief formation that reliably (i.e., regularly) leads to true beliefs (for details, see Goldman, 1979). Thus, reliable processes such as perception and memory yield knowledge, but unreliable processes such as wishful thinking and guessing do not yield knowledge, even on those occasions that they result in true beliefs. Reliabilism allows that one can have knowledge even if one does not have any reason to think that the process that causes one’s belief is a reliable one and even if one has no evidence for the proposition believed. What matters is just that the belief be caused in a reliable way.

Although there are differences between the causal theory and the reliabilist theory, for present purposes their differences are far less significant than their similarities. In what follows, I will sometimes refer to all theories relevantly like these two as “causal” theories.

C A crucial difference between good reasons analyses and causal analyses

The discussion so far suggests that there are two distinct categories of theory about knowledge. One category includes the good reasons theories and the other includes the causal theories. A preliminary way to see the internalism/externalism debate is as a debate between defenders of good reasons theories (the internalists) and the defenders of causal theories (the externalists). However, as is so often the case, things are far more complicated than this initial formulation suggests. Before we look at why things are more complicated, it will be useful to make more explicit the key difference between good reasons theories and causal theories. This will bring to light something about what counts as an internal factor.

Good reasons theorists initially thought that what needs to be added to true belief in order to get knowledge is evidence or reasons. Reasons can include other beliefs, perceptual experiences, apparent memories, and so on. Consider your belief that there is a maple tree outside your window. According to the good reasons theorists, your reason for this belief might be another belief, say the belief that the tree has leaves of a particular shape. Or your reason might include the way the tree looks – how it appears to you. While we might ordinarily say that your reason for thinking that the tree is a maple is that its leaves are a particular shape, the fact that the leaves are that shape is not part of your evidence. What you are going on in judging the tree to be a maple is your belief that it has leaves of particular shape, and perhaps ultimately you are going on how the tree looks to you (your perceptual experience). These are internal, mental states you are in. As another example, suppose you believe that it is warm outside now. Your reason for this belief is your feeling of warmth. And this is an internal factor, while the actual outside temperature is not an internal factor. A person's reasons for a belief are things that the person can, at least in the typical case, describe to someone else and cite in support of the belief. In some cases, however, people may have reasons that they cannot describe, perhaps because they lack the proper vocabulary.

The idea, then, is that a person's reasons are the things the person has to go on in forming beliefs, and they will include how things look and seem to the person, the person's apparent memories, and the person's other beliefs. These are mental things. And internalism, at least as it will be construed here, is the idea that these mental things determine certain crucial epistemic facts. As noted earlier, this implies that if two people are alike with respect to the mental factors, then they must also be alike in the relevant epistemic ways. Thus, if an internalist holds that justification is an internal epistemic matter, then that internalist is committed to the view that if two people are mentally alike in all ways that bear on the justification of a particular proposition, then either they are both justified in believing that proposition or they are both not justified in believing that proposition. Equivalently, if they differ with respect to justification, then there must be an internal or mental difference.

Notably, the things that the causal theorists emphasized in constructing their theories are external factors. The fact that one's belief is causally connected in some particular way to some state of the world is not a fact internal to one's mind. Nor is it, by itself, an evidential fact. (More will be said about this topic in section IIIB.) Of course, one can have evidence about causal connections, but that is different. Thus, the good reasons theorists said that knowledge required true belief plus the right kind of internal factors, namely good reasons or evidence to support the belief. The causal theorists said that knowledge required true belief plus the right kind of external factor, namely a causal connection of the right sort.

As a first pass, then, we might take internalism to include theories relevantly like the good reasons theories and externalism to be theories relevantly like the causal theories. One way to state the idea is that internalism is the idea that knowledge requires justification whereas externalism is that the idea that it requires the right kind of causal connection rather than justification. But framing the issue this way is adequate only if we are careful about several key terms and we set aside some issues that some philosophers have seen as central. In the next section, I turn to those matters.

II Clarifications

A Knowledge and justification

The discussion so far conforms to one way philosophers use the word “justified.” This usage necessarily associates justification with good reasons and evidence. On this way of speaking, a belief is justified only if one has good reasons for it. Accordingly, when the word is used this way, we can say that causal theorists hold that knowledge requires true belief plus the right kind of causal connection, and not true belief plus justification. The preliminary statement of the internalism/externalism theory works out well on this usage (setting aside complications to be described later).

However, it is possible for causal theorists to use the word “justified” in a different way. They can agree with traditionalists that knowledge requires justified true belief, but differ with traditionalists over what is required for justification. These causal theorists would say that a belief is justified when it is caused in the right way and deny that justification always requires good evidence. When the word is used in this way, the dispute between internalists and externalists cannot be formulated as dispute about whether knowledge requires justification, since all sides agree that it does.

There is no point in arguing about the proper use of the word “justified.” What does matter is to be clear about which way we are using the word here. It will be somewhat more convenient in what follows to follow this first way of describing things. On this alternative, then, the controversy between internalists and externalists is, in its initial formulation, over whether knowledge requires justification.

B Justification and well-foundedness

Supporters of versions of the traditional good reasons views acknowledge a distinction between merely having good reasons for believing a proposition and believing that proposition on the basis of those good reasons. The difference emerges in cases in which a person has good reasons for believing some proposition, but ignores or misevaluates those reasons, yet believes the proposition out of wishful thinking or on the basis of some mistaken inference. This has a moral analogue: doing the right thing for the wrong reasons. In the belief case, we might say that the person believes the right thing on the wrong basis. Some philosophers will say that the person is justified in believing the proposition (since she has good reasons) but that she does not believe it justifiably (since she bases her belief on something other than good reasons). Another way to put the latter point is to say that the belief is not well-founded (see Feldman and Conee, 1985, pp. 23–25). I will use this terminology here.

Internalists will agree that whether a person is justified in believing a proposition is an internal matter. It is less clear what to say about whether a belief’s being well-founded depends only on internal factors. At this point, some lack of clarity about exactly what counts as internalism emerges. Suppose a person is mixed up about what one of his own beliefs is actually based on. He has some good reasons for believing some proposition and he thinks that he believes it on the basis of those reasons. In fact, it is vanity or wishful thinking rather than those reasons that actually causes him to have the belief. This shows that he is justified in believing the proposition but his belief in it is not well-founded. Perhaps this makes well-foundedness not a fully internalist

notion. This depends upon whether the fact that his belief is based on wishful thinking rather than his good reasons counts as a mental (or internal) fact. It is hard to say. This is a fact about the causal relations among one's mental states. Perhaps no causal facts count as internal. On the other hand, one might also say that these causal facts are about mental causes, and thus are internal to the person's mind. It is unclear, therefore, whether internalists will count well-foundedness as an internal matter.

Rather than attempt to resolve what really counts as internal, it may be better for present purposes to grant that well-foundedness is not a purely internal concept. Still, a belief is well-founded only if the believer is justified in believing the proposition. And an internalist will argue that whether a person is justified in believing a proposition is determined by reasons or evidence, where these are internal matters in even the most restrictive sense. Thus, their view is that justification is an internal matter, and it is necessary for knowledge. Externalists, presumably, will deny this. They will say that justification, so construed, is not necessary for knowledge.

C The Gettier problem

Internalists have long recognized that a well-founded true belief can fall short of knowledge. That is, they realized that one can base a true belief on excellent reasons yet it can still be not a case of knowledge. It is easiest to see how this can happen by noticing first that well-founded beliefs can be false. Consider anything that you do know. Take as an example the proposition that your neighbor owns a Ford. You have excellent reasons for this: you have seen him driving a Ford, he often talks about the Ford he owns, he has proudly shown you a title to the car with his name on it, and so on. You know that he owns a Ford. Now, imagine an odd but possible alternative example. You have exactly the same reasons, with exactly the same reasons to trust your previously honest neighbor. But, in this alternative case, he is faking Ford ownership. The car he drives is owned by his wealthy uncle, the ownership papers he has are forged, and so on. He has pulled off an elaborate hoax. We need not speculate on his motives for this. In this second example, you have a well-founded belief that he owns a Ford, but your belief is false. You don't know that he owns a Ford, and the traditional analysis gets this case right. It implies that you don't know because your belief is not true. Of course, you did have knowledge in the original, normal, case. And the traditional analysis gets this example right as well.

There is, however, an even odder variation on the case. This case makes life difficult, but interesting, for epistemologists. In this odder case, your belief that your neighbor owns a Ford is true, but not for the reasons you have. Suppose that, in addition to the Ford your neighbor drives, his wealthy uncle also has recently bought a Ford and had the ownership of this other Ford assigned to your neighbor. So it is true that he is a Ford owner, but not for the reasons you are aware of. In this situation, you have a well-founded true belief that he is a Ford owner, but you lack knowledge of this fact. Examples like this one are called "Gettier cases," after Edmund Gettier (1963), who first brought them to the attention of epistemologists.

The reason why Gettier cases are important for the present discussion is as follows. Philosophers who are firmly in the good reasons tradition agree that knowledge is not justified (or even well-founded) true belief. They think that there is something else to be added as well. We need not dwell on the details of the possible fourth conditions

such traditionalists will place on knowledge. It will suffice to examine briefly one line of thought. One idea is that knowledge requires, in addition to justified true belief, that there be no *defeater* for one's belief. According to one prominent account, a defeater is a true proposition such that, if the person were justified in believing it, then the person would not be justified in believing the target proposition. The idea is that this is a missing truth that ruins the person's knowledge. So, on this view, knowledge is undefeated well-founded true belief.

Whether a belief is defeated or not is by all accounts something external to the mind of the believer. That is, things could be exactly the same from the perspective of two believers, yet one might be subject to a defeater and the other not. Two people could be internally alike, yet one is the victim of a Gettier case and the other is not. The examples about your neighbor's Ford illustrate this. Thus, good reasons theorists do not think that what must be added to true belief in order to get knowledge is merely something internal. Instead, what they think is that what must be added is good reasons (which are internal) plus two other conditions: that the belief is based on those good reasons, and that there be no defeaters. The latter is definitely an external condition and the former may be, depending upon just how "internal" and "external" are understood.

In modifying the traditional theory in these ways, good reasons theorists are not abandoning their original approach and accepting something like a causal theory. Good reasons remain essential for knowledge. We can, however, continue to interpret the internalism/externalism debate as the debate between the modified good reasons theories on the internalist side and the causal theories on the externalist side. The causal theorists may also have to say something to deal with the Gettier problem. Or, they may think that the causal connection condition replaces the combination of the justification condition and the condition internalists use to deal with the Gettier cases. In any case, the plausible internalist thesis is that justification is an internal epistemic matter, while knowledge requires something external, such as lack of defeat, to deal with the Gettier problem.

D Deontology

Some philosophers identify internalism with the view that justification involves fulfilling obligations or duties. Similarly, some say that justified belief is epistemically responsible belief. For example, Alvin Plantinga (1993, p. 19) has written that one central internalist idea is that "epistemic justification is *deontological* justification. ... All that it requires is that I do my subjective duty, act in such a way that I am blameless." Exactly how one connects this conception of justification to internalism is a matter of some controversy. I will not pursue it here. What is crucial to note is that the view that epistemic facts depend upon internal facts is logically distinct from the view that epistemic facts are in some sense matters of duty (for discussion, see Conee and Feldman, 2001). That is, it is one thing to say that justification is a matter of doing one's duty or doing what one ought to do. It is another thing to say that the relevant epistemic duty is determined by internal facts (such as evidence). It is possible to emphasize either the duty fulfillment aspect of the view or the internal factors aspect of the view when one characterizes what internalism is. The present discussion takes the second approach. In what follows it is not assumed that internalists are committed to a deontological approach.

As I and Earl Conee have argued elsewhere (Conee and Feldman, 2001, section I), any decision about what aspect of the traditional theories to emphasize in characterizing internalism is somewhat arbitrary. Internalism is the name given to the traditional theories. But those theories were stated prior to the names “internalism” and “externalism” being used to designate epistemological theories. Armstrong introduced the term “externalism” to characterize his causal theory and contrast it with the traditional theories. One could think that what was crucial to the traditional theories was the fact that they made good reasons necessary for knowledge. One could think that what was crucial was that they made justification a deontological matter. One could think that one of these aspects – say the deontology – was supposed to support the other. In this essay, I am identifying internalism with good reasons theories and their descendants, and setting aside the deontologism.

III Internalism Vindicated

Before turning to my general defense of internalism, I want to respond to John Greco’s radical claim, in his contribution to this debate, that “all interesting epistemic evaluations are externalist evaluations.”

A Some internalist epistemic evaluations

It is, of course, difficult to argue effectively about which evaluations are “interesting,” since what is of interest to one person may not be of interest to another. In this section I will describe some evaluations that are internalist, at least in the way characterized earlier. I believe that these count as interesting evaluations, but I do not contend that they are among the necessary conditions for knowledge.

1 Consistency. A person’s overall set of beliefs can be evaluated for *consistency* and each individual belief can be evaluated in terms of whether or not it is consistent with the rest of the person’s beliefs. Although it is unclear exactly what counts as an epistemic evaluation, this would seem to be one. And it is internalist: it depends only upon the relations among the person’s beliefs. Of course, whether the beliefs are consistent is a logical matter, and it may be that people are in some cases ignorant of the logical relations among their beliefs. This shows that people may, for example, fail to realize that a new belief is inconsistent with some previously held beliefs. Still, whether the belief is inconsistent with another belief is an internal matter in at least the following sense: two people who are internally alike, in that they have the same mental states, cannot differ with respect to this property. If one has a belief that is inconsistent with the rest of her beliefs, then the other does as well.

2 Identifiable good reasons. A person might at times think about something he believes and ask whether he has any good reason for that belief. The ability to identify (or think of) such a reason is a mental fact about the person. Again, two people who are mentally alike in that they have the same beliefs and are able to identify the same reasons are alike with respect to this epistemic evaluation. And, once again, it is not assumed here that people are infallible with respect to whether they can identify a good reason. They may be mistaken about what counts as a good reason. Nevertheless, whether a person can identify such a reason is an internal mental fact about the

person. Thus, this is a second internalist epistemic evaluation. By and large, the ability to identify a good reason one has will coincide with the more familiar property of being able to state a good reason. But these two evaluations will diverge in rare cases in which a person has a reason but, somehow, is unable to say what it is. In assessing someone else's beliefs, we may often be interested in whether they can state reasons for them. But it is difficult to see why the ability to identify such reasons should be ruled out as "not interesting." We might wonder, for example, whether a person can think of any good reasons for some unusual beliefs, even if we know that the person is unable to communicate those reasons.

There are, then, some internal epistemic evaluations that are of interest. This refutes Greco's extreme thesis. But it leaves open the question of whether there are any internal evaluations that are directly relevant to knowledge. I turn next to that.

B Knowledge without reasons

Some ordinary uses of the word "knows" clearly fit an externalist picture. We sometimes say such things as "The thermostat knows that it is time to turn on the furnace" or "The plants know that winter is over." Presumably, the truth about the thermostat and the plants involves only facts about the reliability of their responses to certain environmental factors. These uses of "knows" do not imply that the thermostat and the plants have any true beliefs, nor do they imply that they have good reasons for believing anything. One might liken ordinary attributions of knowledge to what is said in these cases, and affirm externalism as a result.

Arguably, however, these uses of "knows" are non-literal (see Conee and Feldman, 2004). We often use psychological terms in ascribing properties to things that do not actually have psychological states. We might say of a car that only runs on premium gas that it "does not like standard grades of gas" even though cars are not capable of literally liking or disliking anything. What we say is just a colorful way of saying that the car does not run well on standard grades of gas. Similarly, then, it may be that the knowledge attributions to the thermostat and the plants are not literally true. The thermostat does not *know* that it is time to turn on the furnace. It just is set up to turn on the furnace at the desired time. Similarly, the plants do not *know* that winter is over. They just come out of winter dormancy when winter is over. If this is right, then these uses of "knows" provide no support for externalism as we are currently understanding it.

Perhaps some epistemologists think that the uses of "knows" currently under discussion are literal and that they refute the traditional analysis. Perhaps they think that this resolves the debate in favor of externalism. However, at most it shows that not all kinds of knowledge require reasons, and perhaps that there is no particular kind of internal state that is necessary for all knowledge. (See Ernest Sosa's, 1997, discussion of animal knowledge.) This leaves open the possibility that we can plausibly divide all cases of knowledge into two kinds. One kind includes only the cases that are mere regular reactions to environmental stimuli. The other kind might be termed "discursive knowledge" or "reflective knowledge." It does seem to require good reasons, and thus to support a kind of internalism. So, it is open to internalists to argue that there is a kind of knowledge for which there is an interesting internalist necessary condition.

Finally, even if one does not say that there is a kind of knowledge that requires the possession of reasons, it is difficult to see why having reasons is not a matter of considerable epistemological interest in its own right. Arguably, this is what the traditional debate about skepticism has concerned (see Conee and Feldman, 2004, for discussion). That is, skeptics have asserted, and their critics have denied, that we have good reasons to believe many of the things we ordinarily believe. If we concede knowledge to the externalists, then skepticism is best transformed into a debate directly about the quality of our reasons. Since skepticism, so transformed, is interesting, this epistemic evaluation continues to be of interest.

C Knowledge with reasons

While discussing an argument for internalism based on deontological considerations, Greco points out that “we can make a distinction between (a) *having* good reasons for what one believes, and (b) believing *for* good reasons” (p. 329). This is the same as the distinction drawn earlier between a proposition being justified for a person and the person having a well-founded belief in that proposition. Greco argues that for a belief to be responsibly formed, and thus justified according to a deontological view, it must be that it comes about in the right way. That is, it must be held for good reasons. And this, he assumes, is not an internal matter. I will grant that this is correct.

Greco makes similar claims about a view closer to the non-deontological internalism defended here. He considers the idea that two believers who are internally alike are equally well justified. He replies that

The considerations about epistemic responsibility above suffice to counter this line of reasoning, however. The problem is that two believers might be alike internally, yet different regarding the causal genesis of their beliefs. And once again, etiology matters. Suppose that two persons arrive at the same perspective, but that one does so in a way that is epistemically responsible, whereas the other does [... not.] The two persons will not be alike with respect to epistemic justification, even though they share the same internal perspective. (p. 330)

Here, “epistemic justification” must be used to mean “well-foundedness.” Given that, the conclusion can be granted. However, this fails to show that the two believers will not be alike with respect to having good reasons, the epistemic evaluation to which Greco himself called our attention. Indeed, it seems clear that if they are internally alike, then they are alike with respect to having good reasons, and thus internally alike with respect to this epistemic evaluation. Furthermore, having good reasons is necessary for knowledge. Hence, it follows that there is an internal necessary condition for knowledge – having good reasons, or justification.

This does not quite clinch the case for internalism, if the issue is whether there is an “important” or “interesting” internal necessary condition for knowledge. Since Greco himself calls our attention to the property of having good reasons, one might conclude that he thinks that it is interesting. However, I can think of two reasons for denying that justification is an interesting or important necessary condition for knowledge. Greco explicitly discusses only one of them.

Consider first the reason Greco does not discuss. An externalist might argue as follows. Consider two rival analyses of knowledge:

- Analysis 1: Knowledge is undefeated justified true belief.
Analysis 2: Knowledge is undefeated well-founded true belief.

Philosophers who support anything along the lines of the traditional analysis will prefer Analysis 2 to Analysis 1. Thus, the interesting necessary condition for knowledge, the one that shows up in the best analysis, is well-foundedness rather than justification. It follows that well-foundedness, rather than justification, is the interesting epistemic evaluation in this vicinity, and it is an externalist evaluation. Justification, while internal, is not interesting.

This argument fails. Not all interesting necessary conditions for knowledge show up in a properly spelled out statement of the conditions for knowledge. Note that having good reasons, or justification, is necessary for well-foundedness. So having good reasons is necessary for knowledge. From the fact that justification is not listed as a separate or independent element in Analysis 2, it does not follow that justification is not an interesting genuinely necessary condition for knowledge. What shows up as an “independent” necessary condition of knowledge depends upon inconsequential details about how we happen to write out our analysis of knowledge. Here is a way to rewrite Analysis 2:

- Analysis 3: Knowledge is undefeated justified true belief in which the belief is based on justifying reasons.

Given that well-founded beliefs are justified beliefs that are based on the justifying reasons, Analysis 2 and Analysis 3 are equivalent. Analysis 2 does not explicitly include a specifically internalist component. But this is because it makes use of well-foundedness. Analysis 3 spells out what this element depends on, and in doing so it does make use of an explicitly internalist element. Thus, the (presumed) adequacy of Analysis 2 does not show that knowledge does not have an important necessary condition that is internalist. It only shows that there is a way to write down this analysis that does not make explicit appeal to this element.

There is not the slightest reason to think that justification (or good reasons) is not “interesting” or not “important” just because we can restate the analysis of knowledge in terms that mask its presence. It is not that Analysis 2 makes having good reasons not necessary for knowledge. It just does not make that factor explicit. It puts it inside the “well-foundedness” requirement.

Greco presents a different reason for thinking that justification is not interesting in section 3 of his essay. He argues that all interesting epistemic evaluations are either objective or subjective. The objective evaluations have to do with truth or “fit with the world” and are not internal. The subjective evaluations are the ones that have the best chance of being internalist. But he argues that all such evaluations involve responsibility and this involves etiology and is therefore external as well. However, as I pointed out earlier, an evaluation in terms of whether one has a good reason – what I have been calling “justification” – does not depend upon either responsibility or etiology. I see no reason to reject this evaluation as “uninteresting.” Consider some beliefs of importance, say religious beliefs or beliefs that figure centrally in your views about morality. A question you might ask yourself is whether you have any good reason to think that these beliefs are true. The claim that the evaluation that results is uninteresting is without merit.

D Externalist accounts of good reasons

The argument of the preceding section depends upon an assumption about reasons that should be made explicit. It is that what counts as a good reason is not an external matter. This is not to say that everyone always knows what counts as a good reason. It is just to say that external factors cannot make one thing a good reason in one case but not in another. External factors cannot make it the case that two believers who are internally alike differ in what they have good reasons to believe. Some externalists, such as Christopher Hill (1999), reject this assumption.

We can use an example Greco describes to develop this point. Suppose that a person learns the history of the country from a source that is unreliable but which he has every reason to trust. Greco says that there is a clear sense in which this person's historical beliefs are "subjectively appropriate." He goes on to argue that subjective appropriateness depends upon past factors. But the key thing to notice here is that in this example, the beliefs are subjectively appropriate when they are formed (since the person has every reason to trust his source) and they remain that way later on. This is just what an internalist would say about the case.

To reject the idea that this internalist evaluation is an "interesting" epistemic evaluation, one would have to say that the belief is not subjectively appropriate, either when formed or later on, simply because the source was not in fact reliable. Now, that is not a judgment Greco is prepared to make. However, it is possible that some externalists will say that the person has a good reason only if the source actually is reliable, no matter what information the person happens to have about its reliability. But this response is surely implausible. There is, of course, a difference between relying on a reliable source that one has reason to trust and relying on an unreliable source that one has equally good reason to trust. Yet there surely is something common to the cases, something favorable to be said of the person who does accept the word of both sources. The person who accepts what one such source says and rejects what the other says might, by chance, hit upon the truth. But such a person would surely not have reason on his side.

E Forgotten evidence examples

Examples in which people forget their original basis for a belief provide a key element of some objections to internalism. I examine them in this section.

Greco describes an example of a person, Maria, who has a clear apparent memory that Dean Martin is Italian and no current reason against the proposition that he is Italian. However, she initially formed this belief irresponsibly and unjustifiably, relying on the testimony of someone she knew to be untrustworthy. She has, however, forgotten that this is her source. Greco's example resembles one put forward by Alvin Goldman (1999, section III) in a critical discussion of internalism.

Greco says that Maria is not blameless in holding the belief that Dean Martin is Italian. His point is that whether a current belief is blameless depends upon its history, not just the believer's current situation. This case comes up as part of his response to an argument for internalism that makes use of premises about blameless belief. I do not endorse any such argument and will not dispute this part of his evaluation of the case.

However, I do want to examine more carefully whether Maria's belief is justified. I will consider two possible responses open to internalists.

Reply 1. Maria's belief is justified, given that her current evidence does clearly support the proposition that Dean Martin is Italian. There is an important assumption that gives this response credibility, and it is worth making this assumption explicit. The assumption is that there is some attitude toward this proposition that it is reasonable for Maria to take (and, more generally, for any person to take toward any proposition the person considers). Suppose Maria considers the proposition that Dean Martin is Italian and wonders what attitude to take toward it. She has a clear memory of learning of this, and has good reason to trust her memory. She has, as the statement of the example makes explicit, no reason to think otherwise. It would be absurd for her to think, in spite of all this, that he is not Italian. So, disbelieving the proposition is clearly not a reasonable option, given the situation she is in. Perhaps a critic thinks that she would be most reasonable to suspend judgment. But this, too, is quite implausible. She has reasons to think he is Italian and no reason to think otherwise. Nothing competes with her reasons in favor. She might appeal to some general skeptical worries – one's memory can always lead one astray – but this is not relevant here. Thus, of the options open to her – believing, disbelieving, suspending judgment – believing is the only sensible option. Her belief is justified after all. This justification is determined by internal factors.

As noted, an assumption behind this argument is that some attitude or other is the reasonable one for her to take. This is a plausible assumption. How could it be that no option is reasonable, that whatever attitude she takes, it would be epistemically bad? Of course, it could be that no attitude will give her knowledge. She might not be in a situation in which she can know the truth of the matter. Still, some attitude or other must be reasonable for her. It is very hard to see how any attitude other than belief could be the one that is epistemically best in her situation. It is true that she made a mistake earlier. And perhaps this mistake would prevent her from having knowledge even if her belief is true. That is, if her belief is true, then, as Conee and I (2001) argue, this is a Gettier case.

I conclude that, whatever we say about blamelessness, Maria is currently justified in believing that Dean Martin is Italian, just as an internalist view suggests. Perhaps some readers will not be convinced. Perhaps they will insist that Maria's belief is a bad one, epistemically speaking. Another response suggests an additional point about the resources open to internalists.

Reply 2. Notice that as Greco describes the case, Maria previously had reason to distrust her source for this belief. Internalism implies that her belief was not justified when she still had this reason. Greco's view is that having this origin makes her belief not justified at the later time, after all information about its bad origin is forgotten. Suppose we go along with this. What exactly this would show about internalism is far from clear. Much turns on just what counts as internalism. At most, what this argument shows is that past internal states matter. So suppose one holds that justification is a matter of one's history of internal states, not just one's current internal states. This presents a puzzling question: is such a view a kind of internalism?

As I see it, there is no definitive answer to this, since what counts as internalism is at least in part a matter of stipulation. But it is worth noting that such a view does not vindicate anything very close to the causal theories. It differs from current state internalism only in that it makes past internal states matter. Reasons still matter. It is just that one's history of reasons matters, as well as one's current reasons.

IV Conclusion

I conclude that the case for internalism in epistemology is very strong. The internalism in question is the view that certain interesting and important epistemic evaluations depend entirely on internal factors, namely reasons or evidence. There are, of course, epistemic evaluations that are not internalist. These include knowledge, being defeated (or undefeated), and, perhaps, well-foundedness. Still, justification, construed along traditional lines, remains an important necessary condition for knowledge. And this suffices to vindicate internalism.

References

- Armstrong, D. (1973) *Belief, Truth and Knowledge*. Cambridge: Cambridge University Press.
- Chisholm, R.M. (1957) *Perceiving: A Philosophical Study*. Ithaca, NY: Cornell University Press.
- Conee, E. and Feldman, R. (2001) Internalism defended. In *Epistemology: Internalism and Externalism*, ed. H. Kornblith. Malden, MA: Blackwell.
- Conee, E. and Feldman, R. (2004) Making sense of skepticism. In *Evidentialism*, pp. 277–306. Oxford: Oxford University Press.
- Feldman, R. and Conee, E. (1985) Evidentialism. *Philosophical Studies* 48: 15–34.
- Gettier, E. (1963) Is justified true belief knowledge? *Analysis* 23: 121–123.
- Goldman, A. (1967) A causal theory of knowing. *Journal of Philosophy* 64: 357–372.
- Goldman, A. (1979) What is justified belief? In *Justification and Knowledge*, ed. G.S. Pappas. Dordrecht: Reidel.
- Goldman, A. (1999) Internalism exposed. *Journal of Philosophy* 96: 271–293.
- Hill, C. (1999) Process reliabilism and Cartesian skepticism. In *Skepticism: A Contemporary Reader*, ed. K. DeRose and T.A. Warfield. Oxford: Oxford University Press.
- Malcolm, N. (1963) Knowledge and belief. In *Knowledge and Certainty: Essays and Lectures*. Englewood Cliffs, NJ: Prentice Hall.
- Plantinga, A. (1993) *Warrant: The Current Debate*. Oxford: Oxford University Press.
- Sosa, E. (1997) Reflective knowledge in the best circles. *Journal of Philosophy* 94: 410–430.